

Tight on budget?

Tight Bounds for r-Fold Approximate Differential Privacy

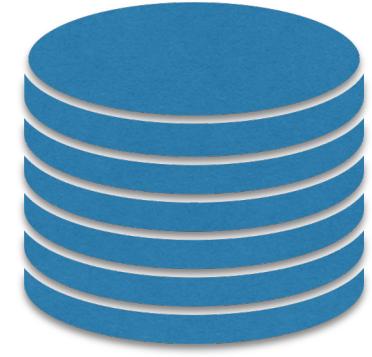
Sebastian Meiser (UCL), Esfandiar Mohammadi (ETH Zürich)

16 October 2018

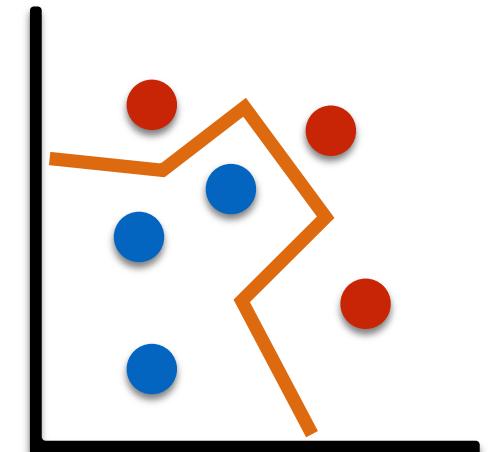
Competing Utility & Privacy Requirements

In many applications competing **utility** and **privacy** requirements

- Statistical queries:
database queries while **protecting entries of individuals**



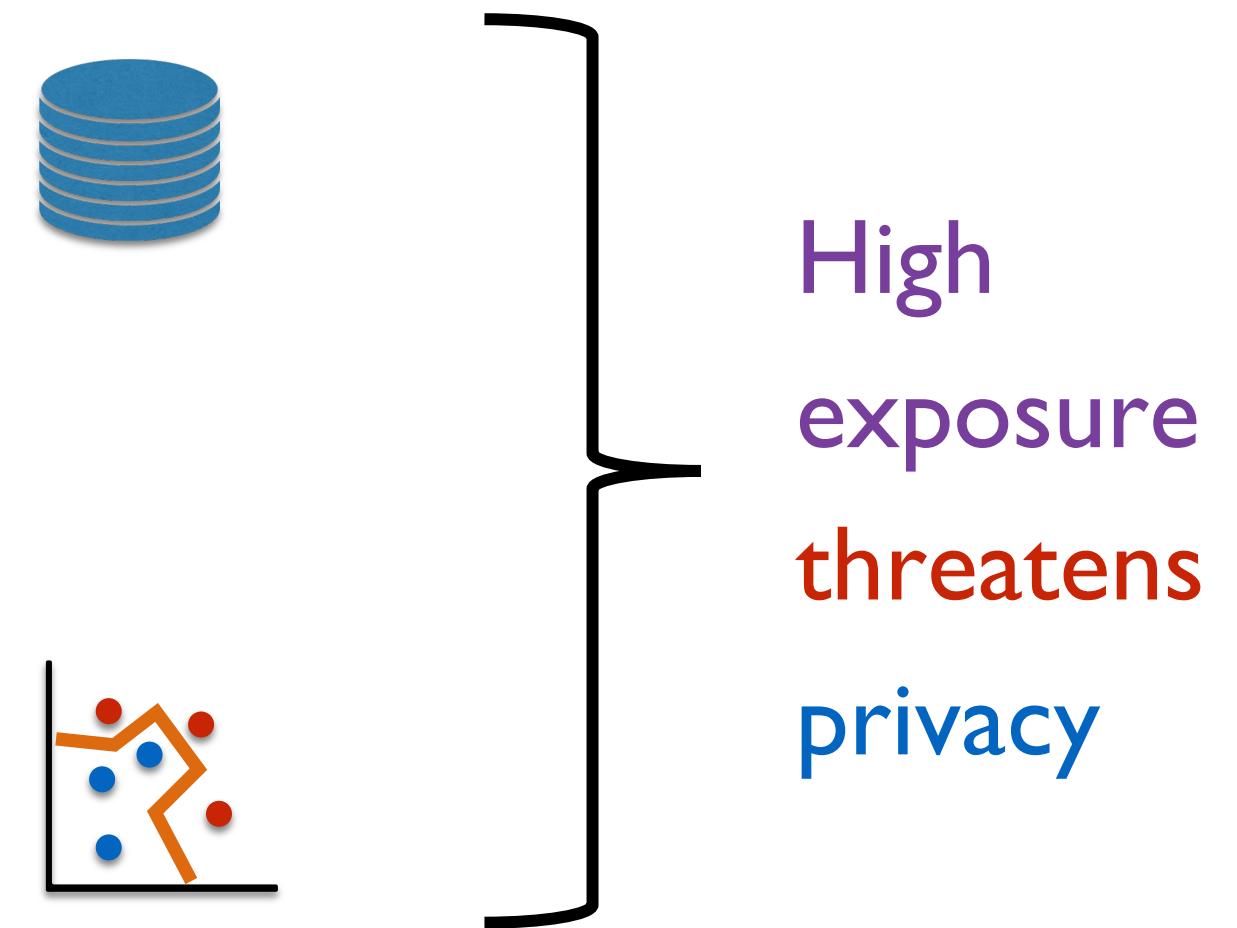
- Private Machine Learning (ML):
queries to ML model while **protecting sensitive training data**



High exposure threatens privacy

In practice: **privacy under continual observation / high exposure**

- Statistical queries:
high nr. of database queries
- Private ML:
high nr. of queries to ML model, or
high nr. of training steps on sensitive data

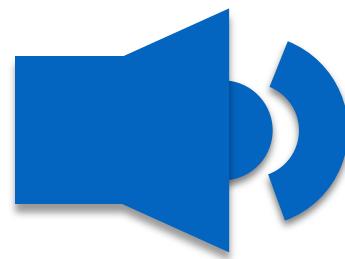


Composition and differential privacy

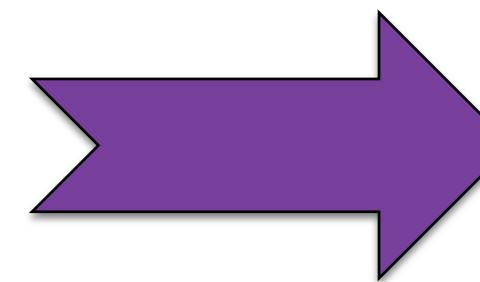
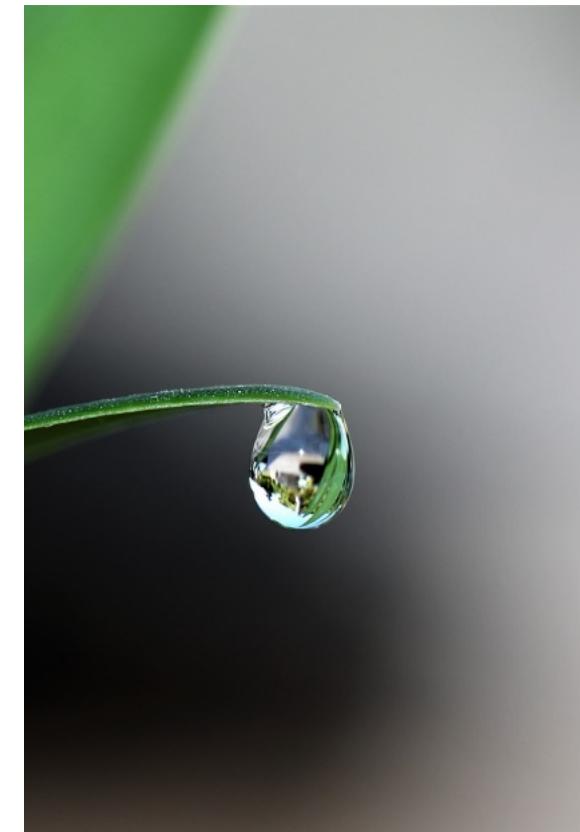
Continual observation / high exposure → leakage **amplifies**

How can we counter that?

Noise



Leakage (1 step)



Leakage (overall)



Composition and differential privacy

Continual observation / high exposure → leakage **amplifies**

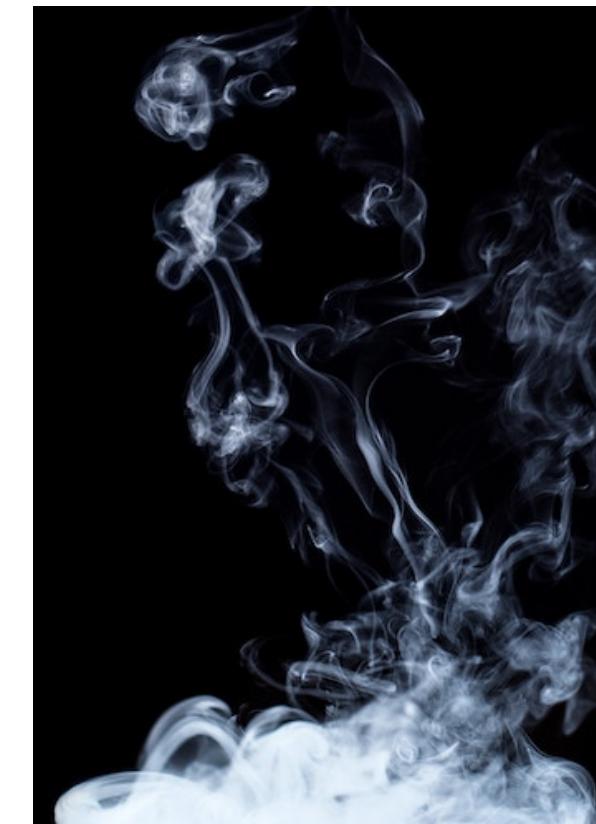
How can we counter that?

- more **noise** to achieve **privacy** under continual observation
- more **noise** renders result of the computation **less accurate**

Noise

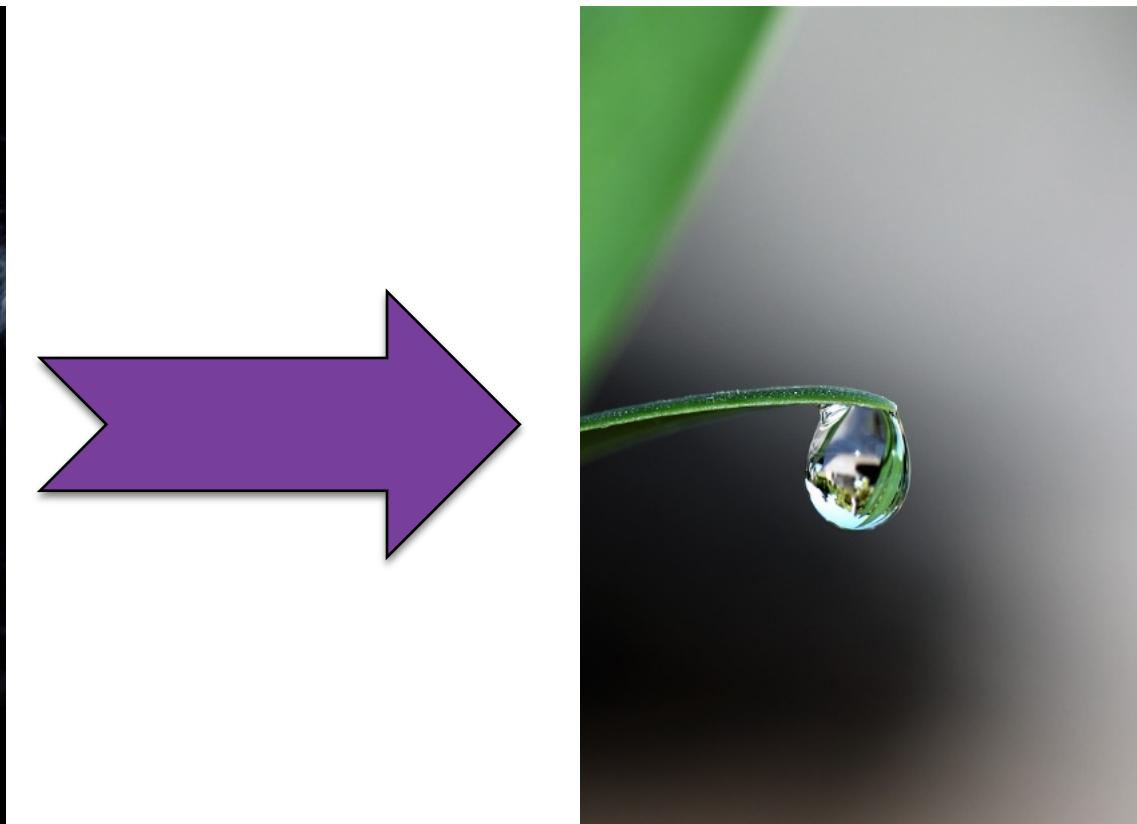


Leakage (1 step)



(water vapor)

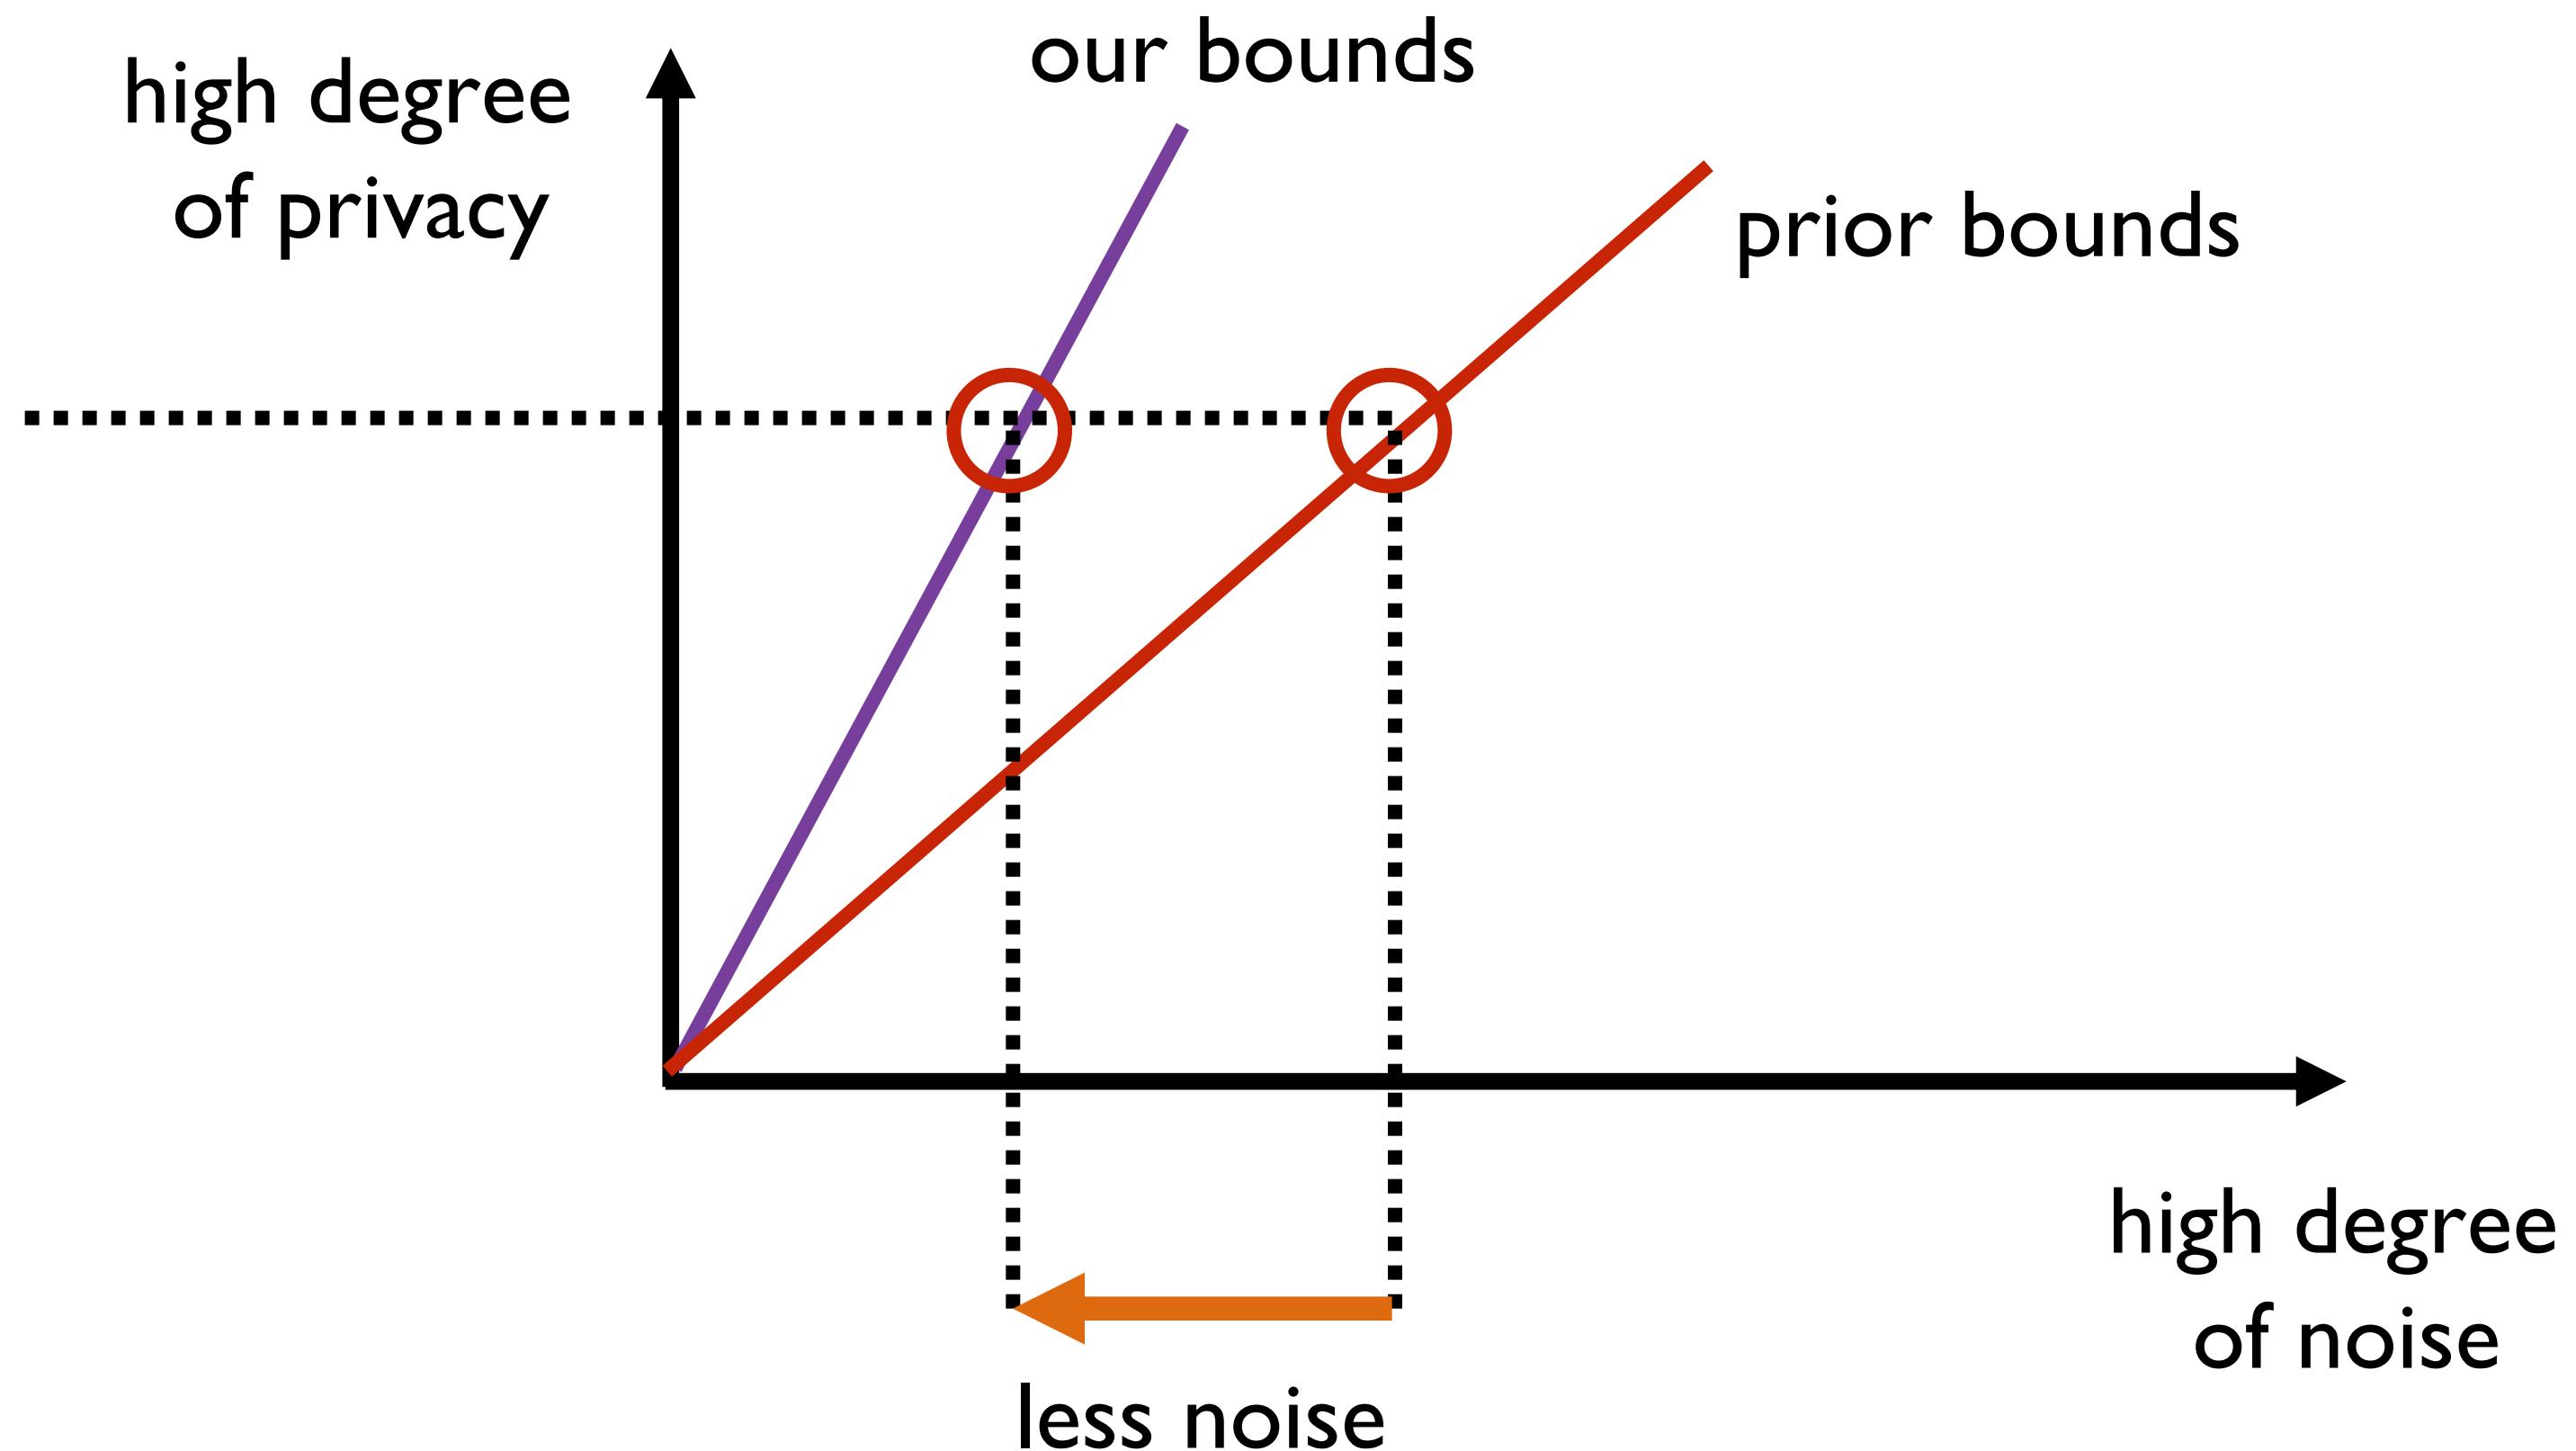
Leakage (overall)



Our work

Previous bounds for approximate differential privacy (ADP)
[widely used privacy notion] under continual observation not tight

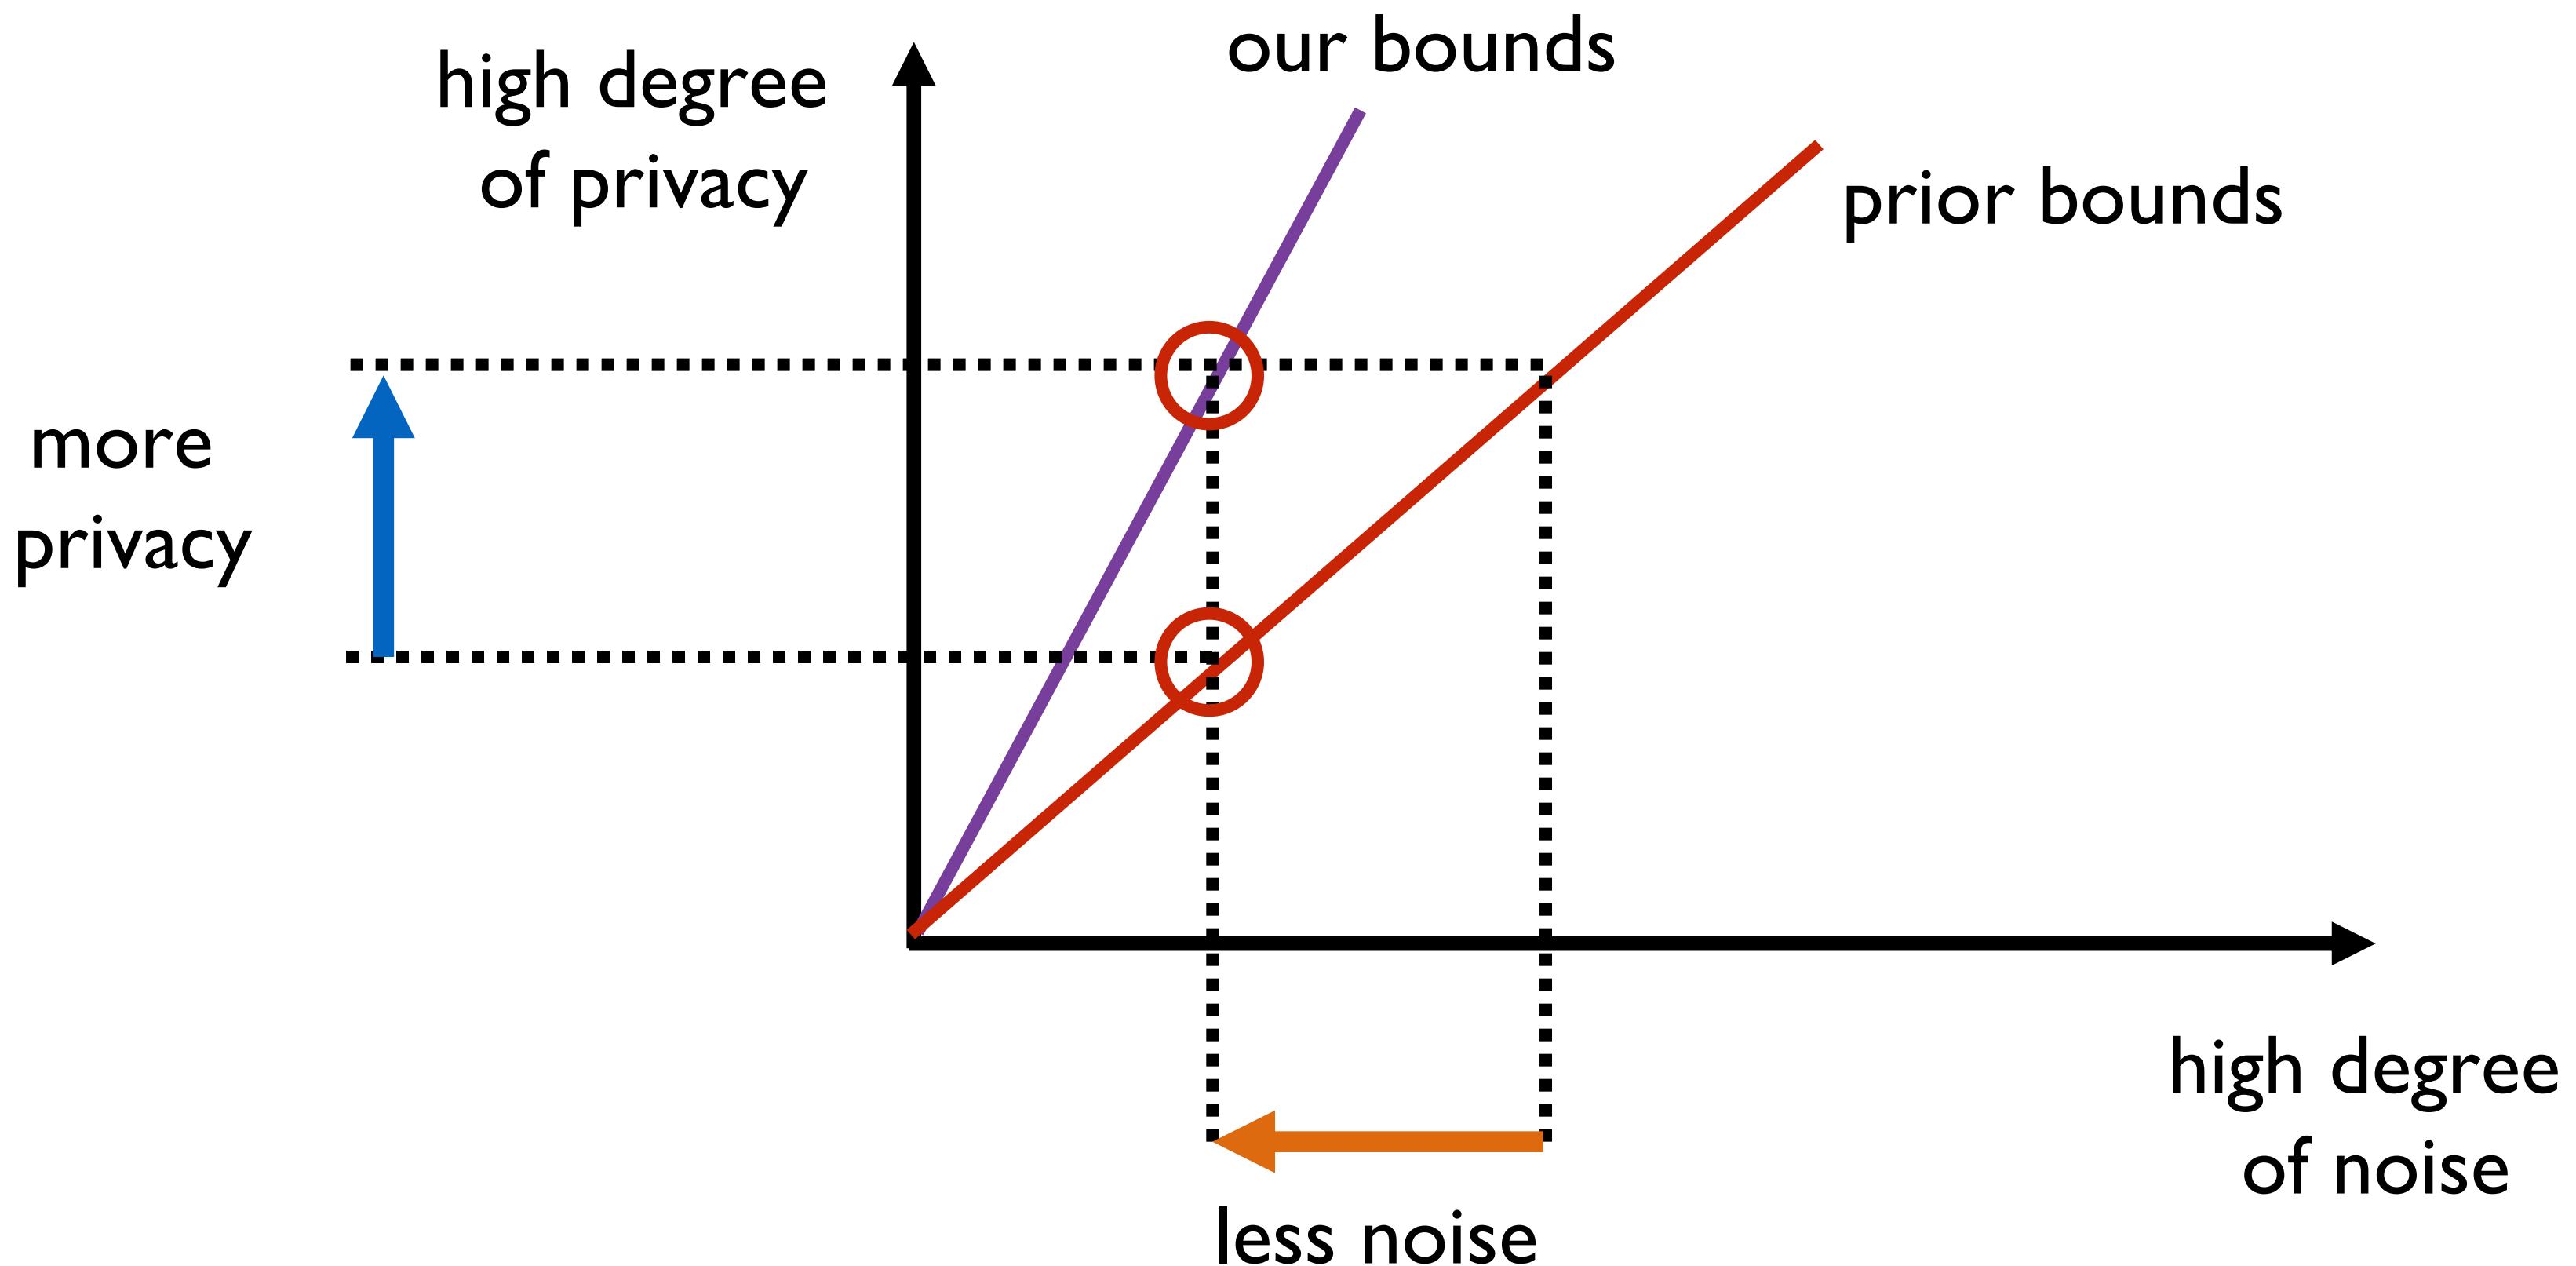
Our work: tight bounds for ADP under continual observation



Our work

Previous bounds for approximate differential privacy (ADP)
[widely used privacy notion] under continual observation not tight

Our work: tight bounds for ADP under continual observation



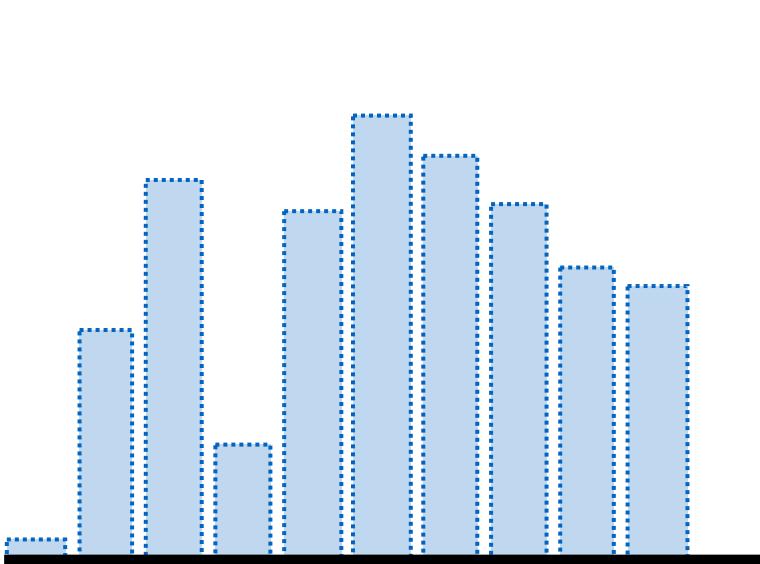
Simple running example: A private mechanism

Consider the following mechanism:

- (Secret) Input: Is Esfandiar or Sebastian watching TV?

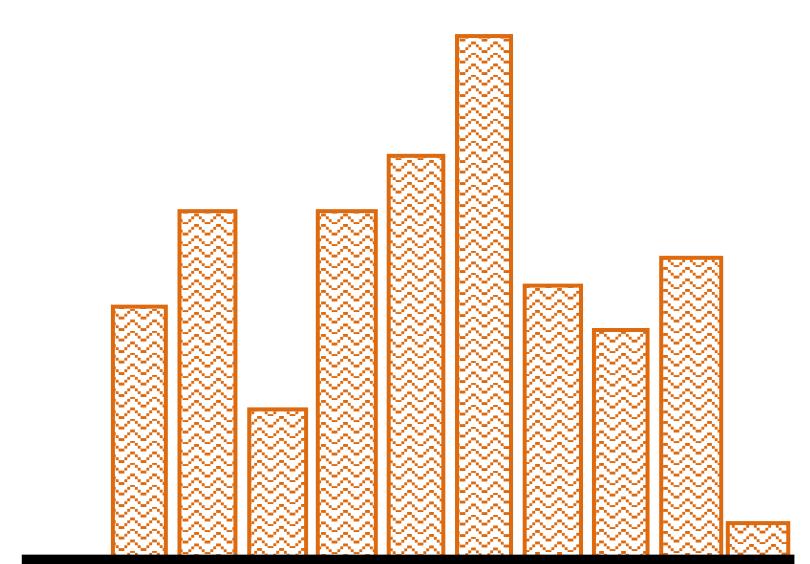
Output: A TV channel.

- The probabilities depend on who is watching.



channels

Sebastian



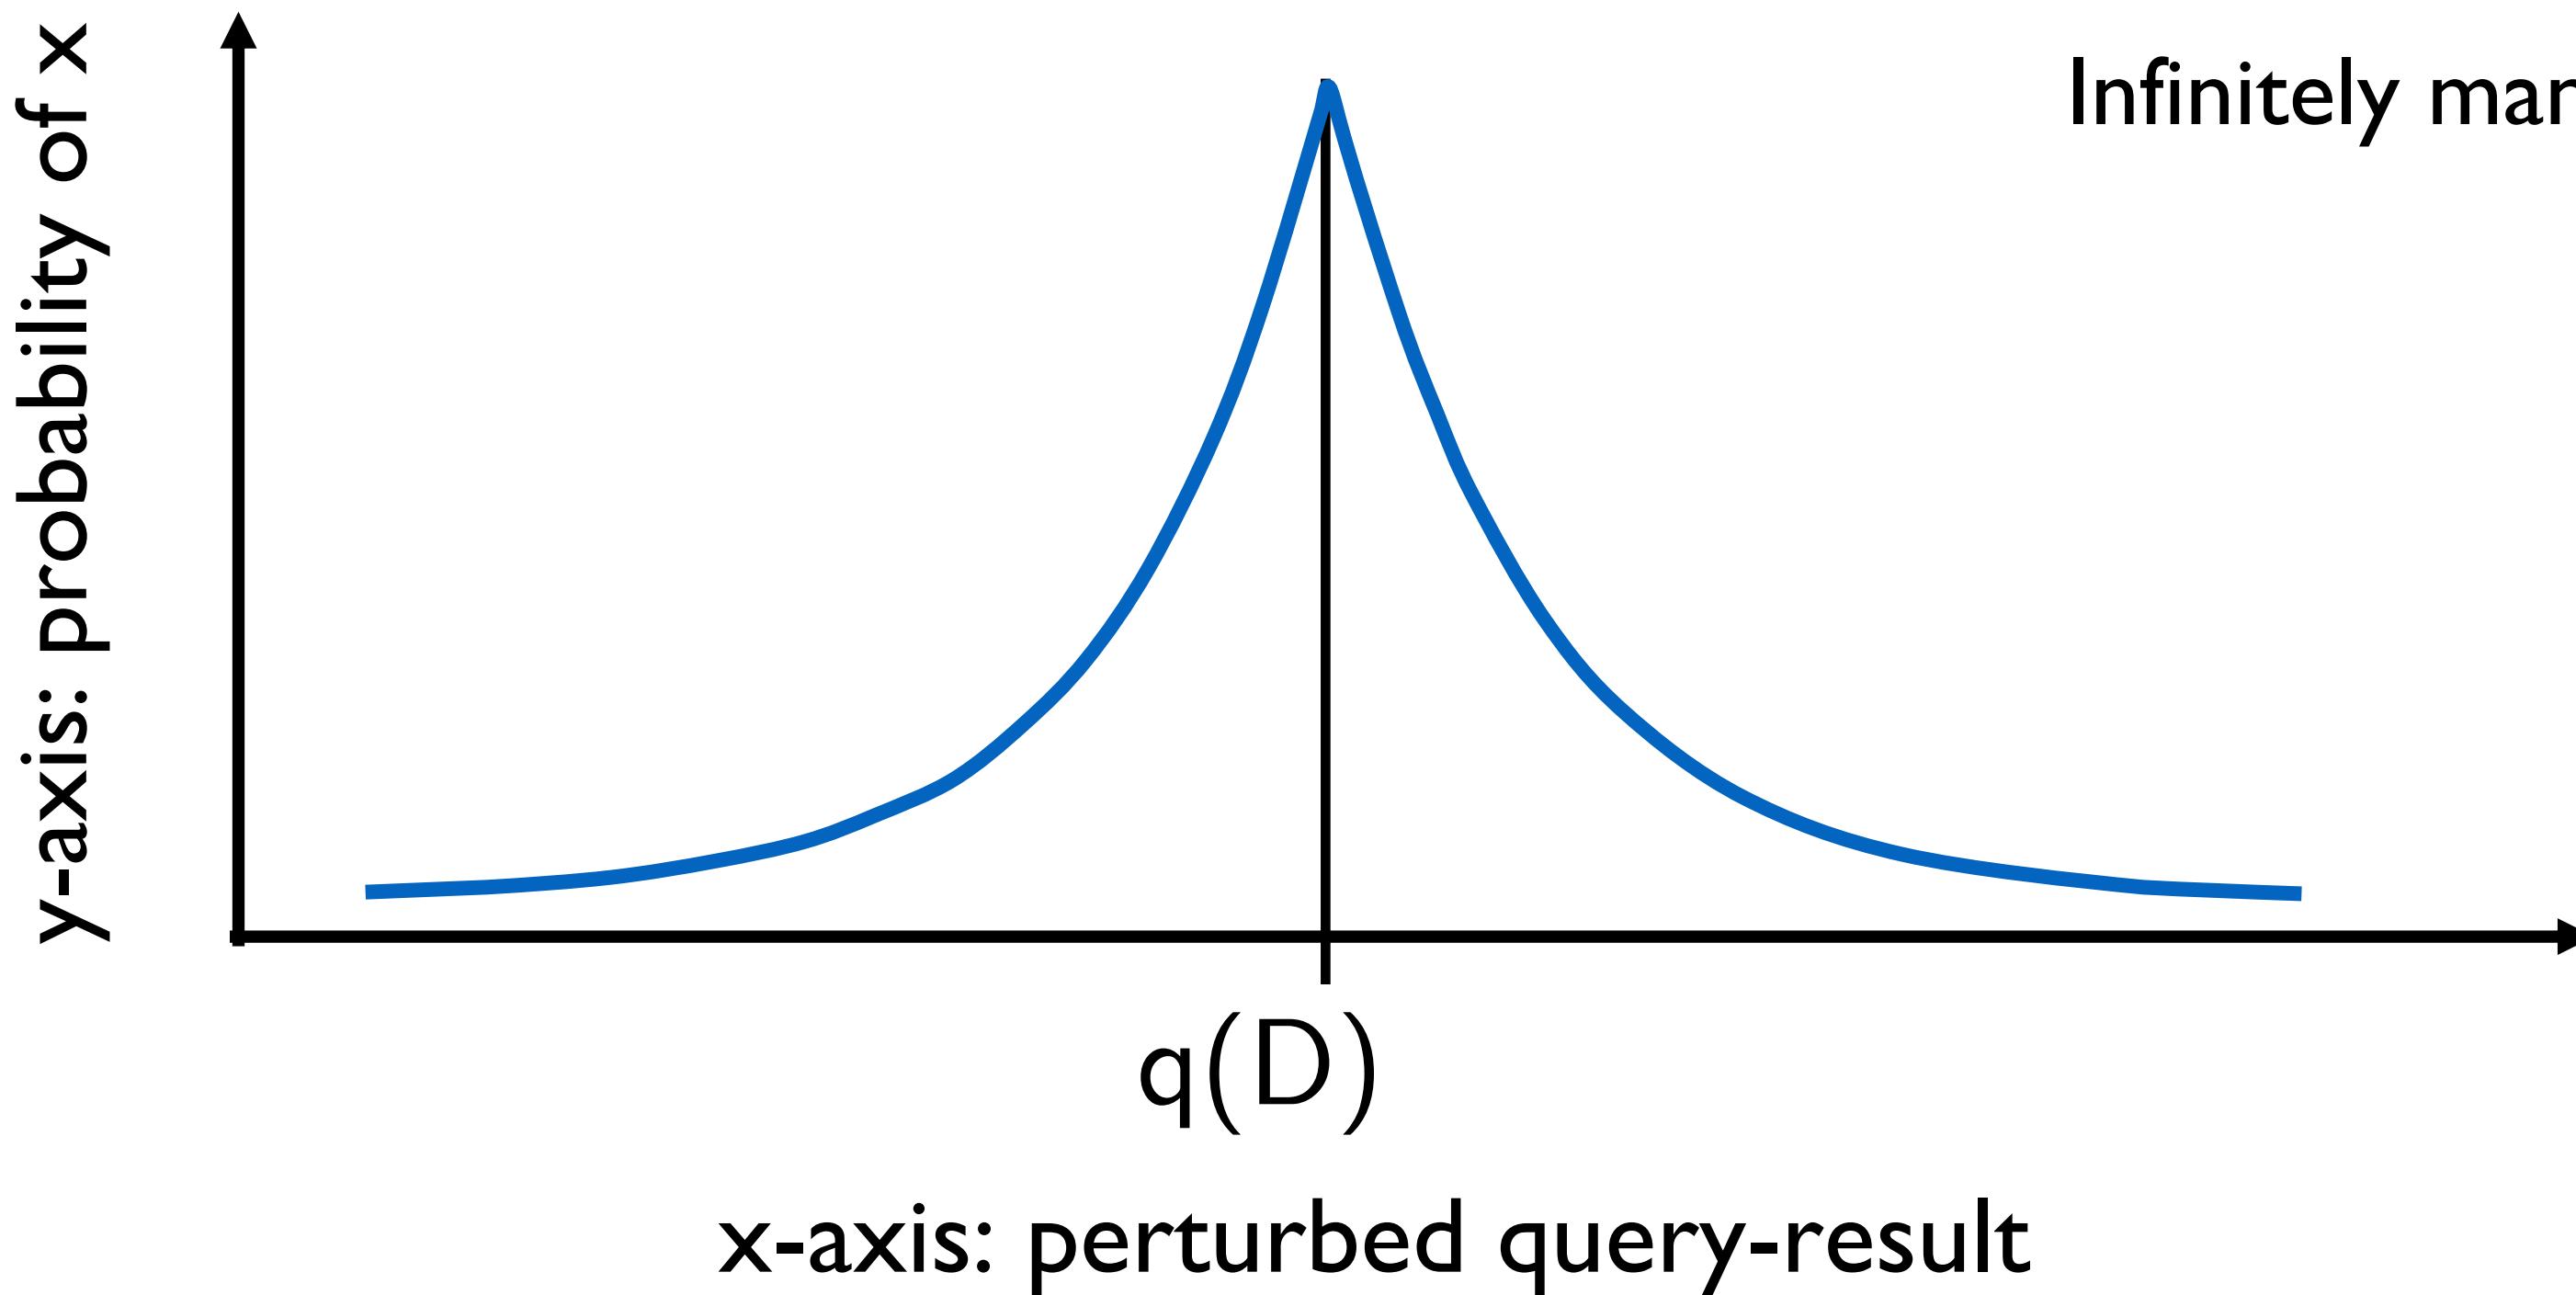
channels

Esfandiar

Why only analyze one pair of distributions?

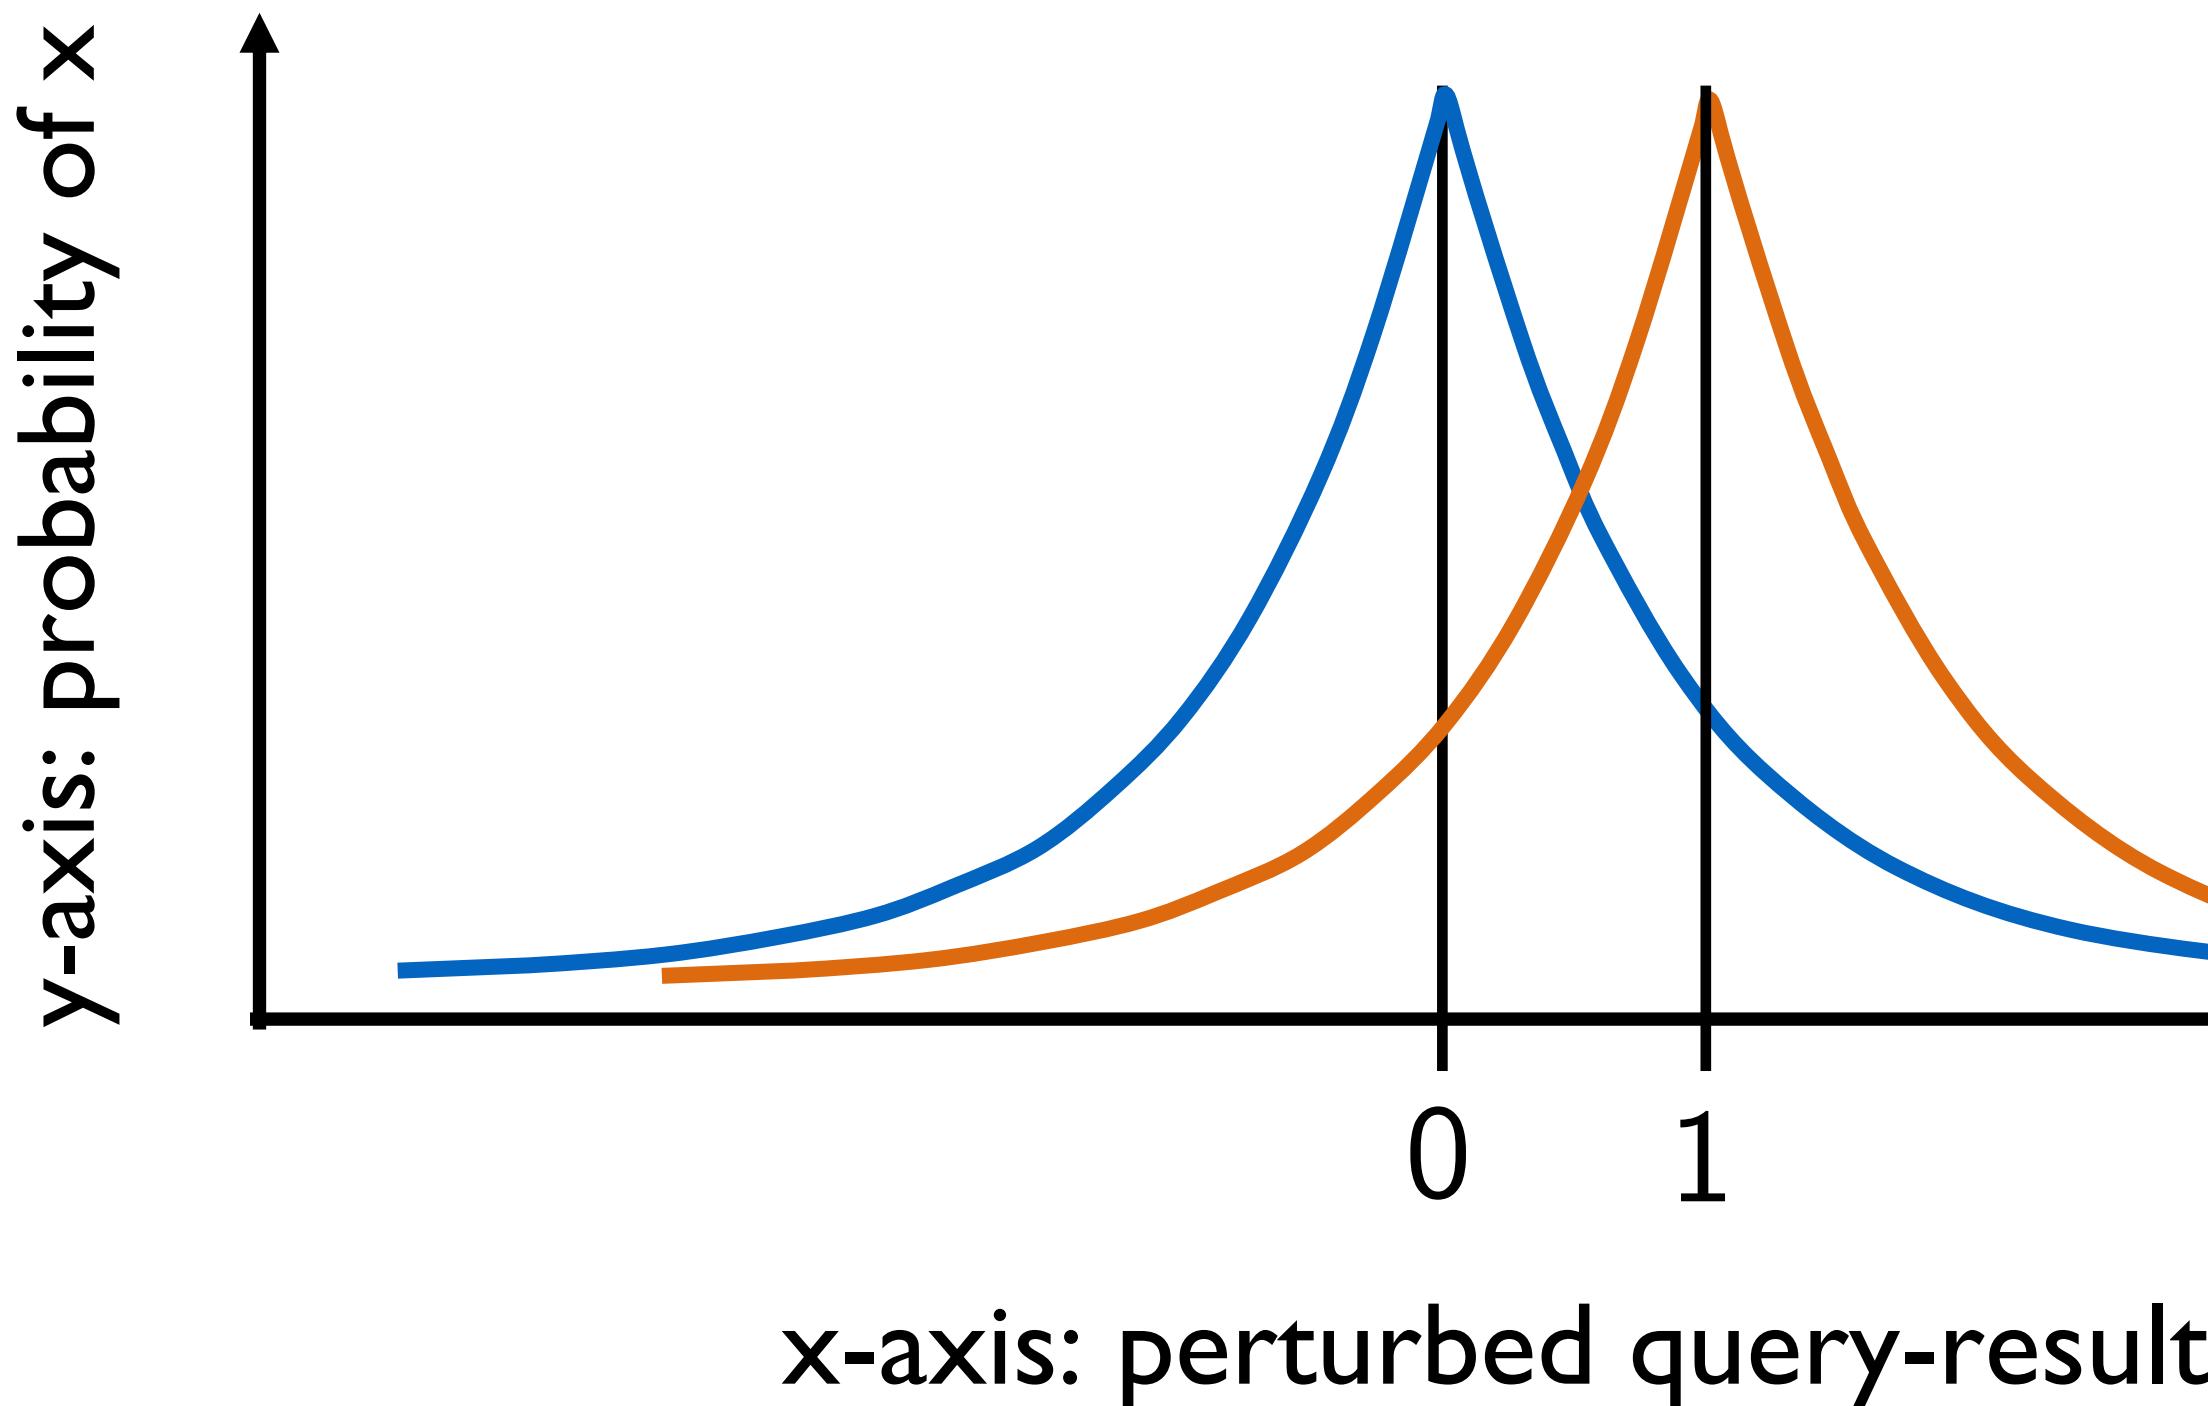
Input: database D , query q

Add λ -Laplace noise
to the query-result $q(D)$
 $\text{Laplace}(q(D), \lambda)$
Infinitely many distributions



Why only analyze one pair of distributions?

Input: database D, query q



Add λ -Laplace noise
to the query-result $q(D)$
 $\text{Laplace}(q(D), \lambda)$

Infinitely many distributions

Worst-case distributions:
 $\text{Laplace}(0, \lambda)$ vs
 $\text{Laplace}(1, \lambda)$
(if sensitivity is 1)

(Approximate) Differential Privacy (for the 4th time?)

Definition: For all pairs x_0, x_1 with $|f(x_0) - f(x_1)| \leq 1$, for all sets S ,

$$\Pr[\text{Seb} \in S] \leq e^\varepsilon \cdot \Pr[\text{Esf} \in S] + \delta$$



ε - δ - Graph

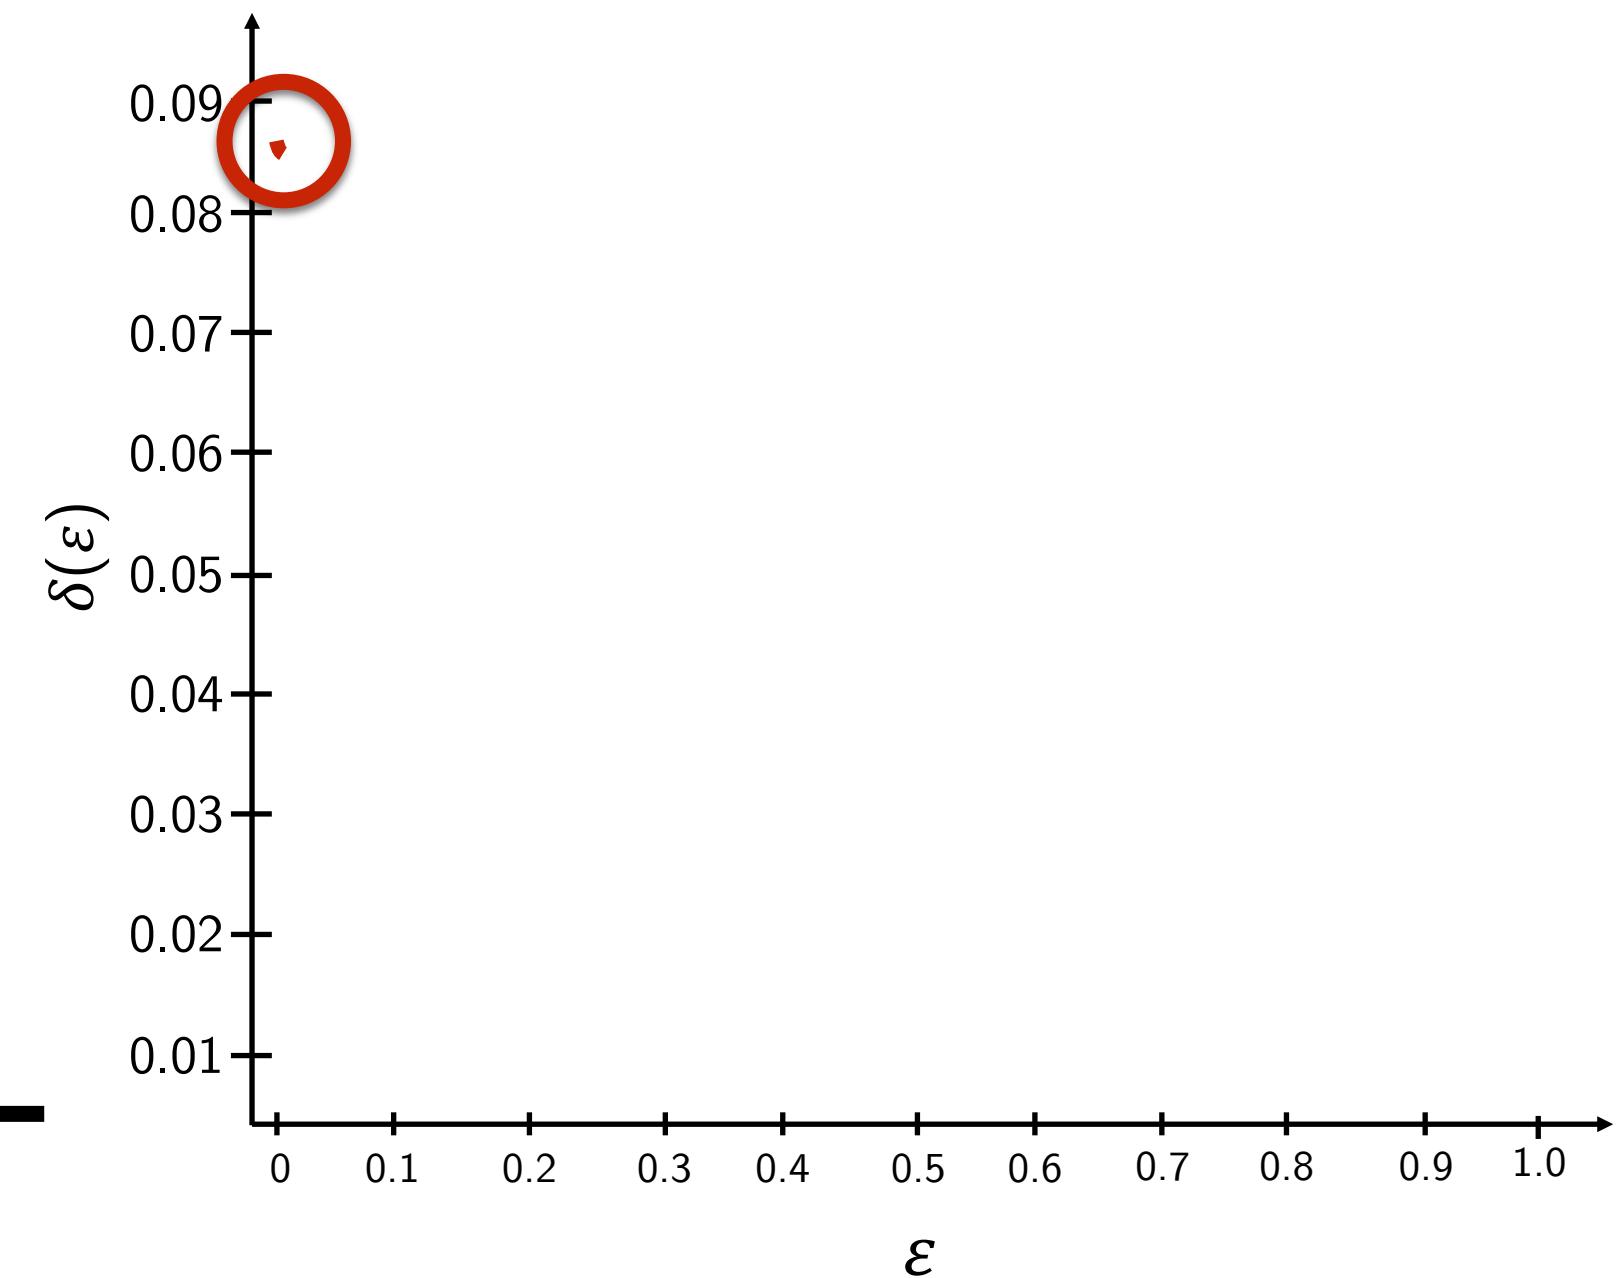
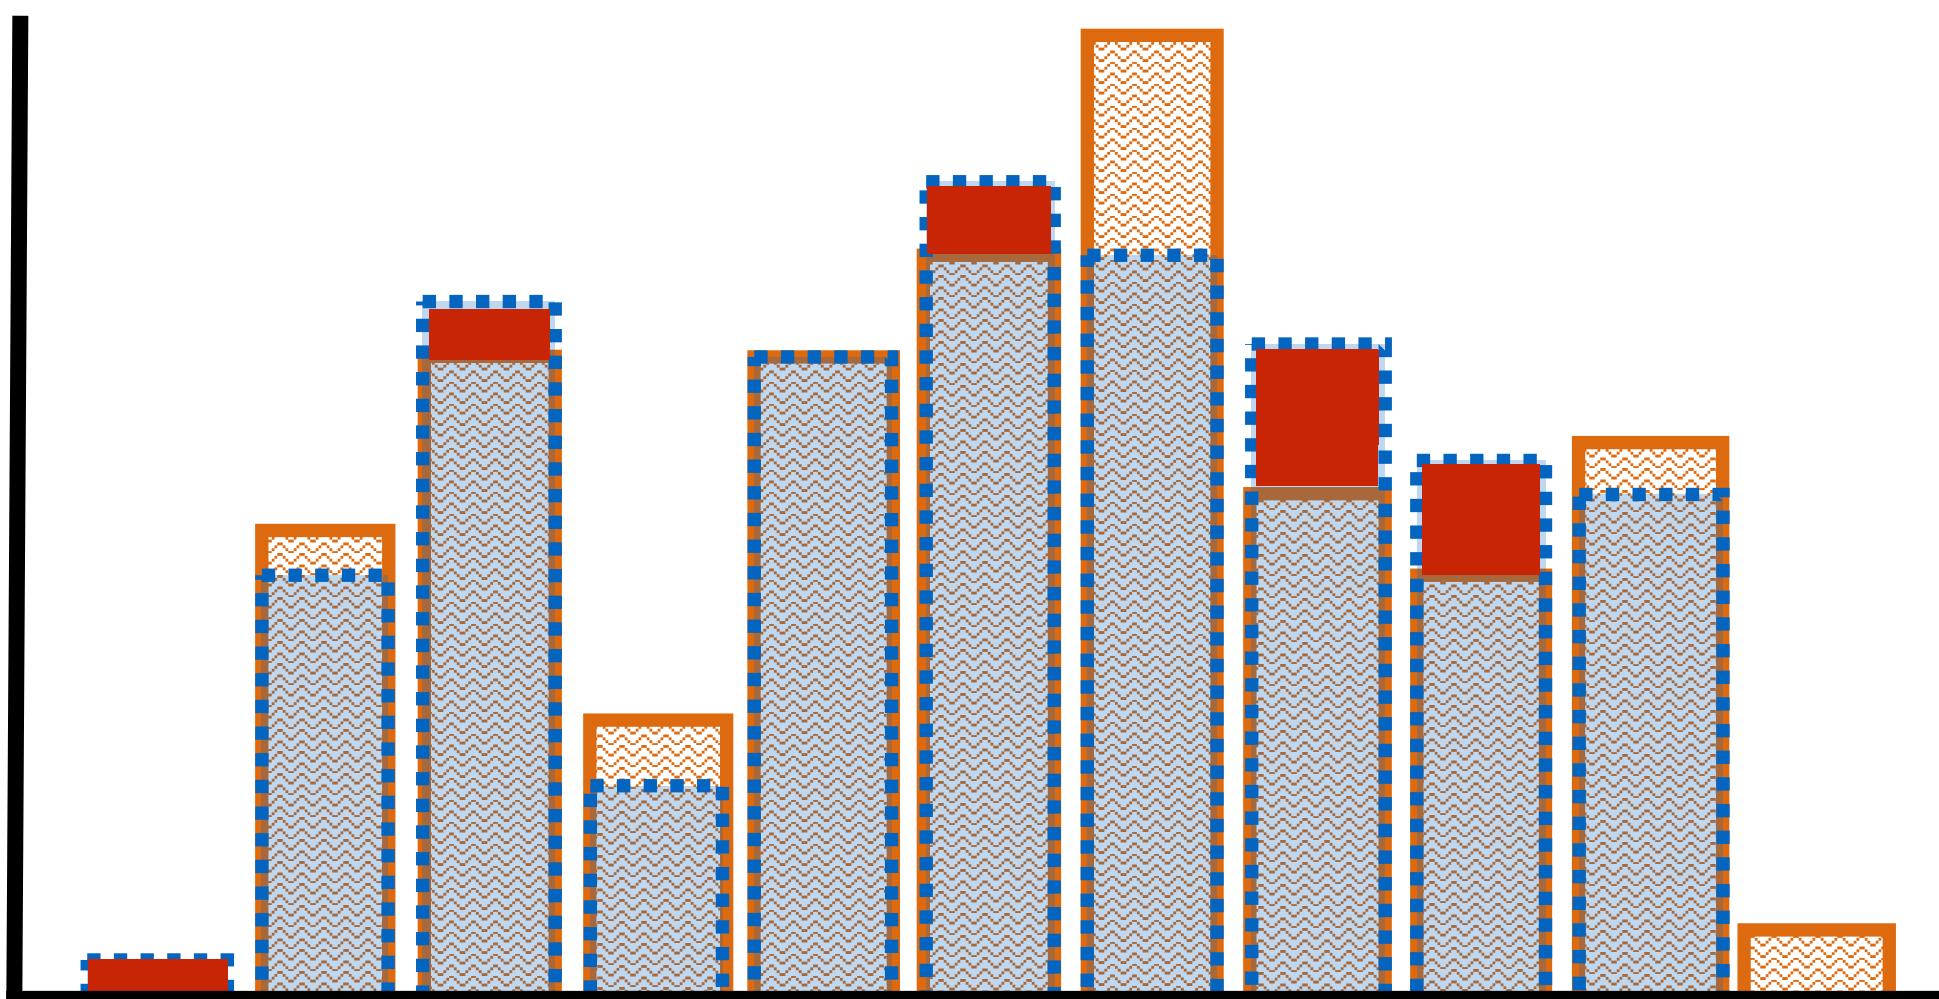
For each ε there is a minimal $\delta(\varepsilon)$

$$\varepsilon = 0$$

$$e^\varepsilon = I$$

A vs $e^\varepsilon \cdot B$

- Sebastian (A)
- Esfandiar (B)



ε - δ - Graph

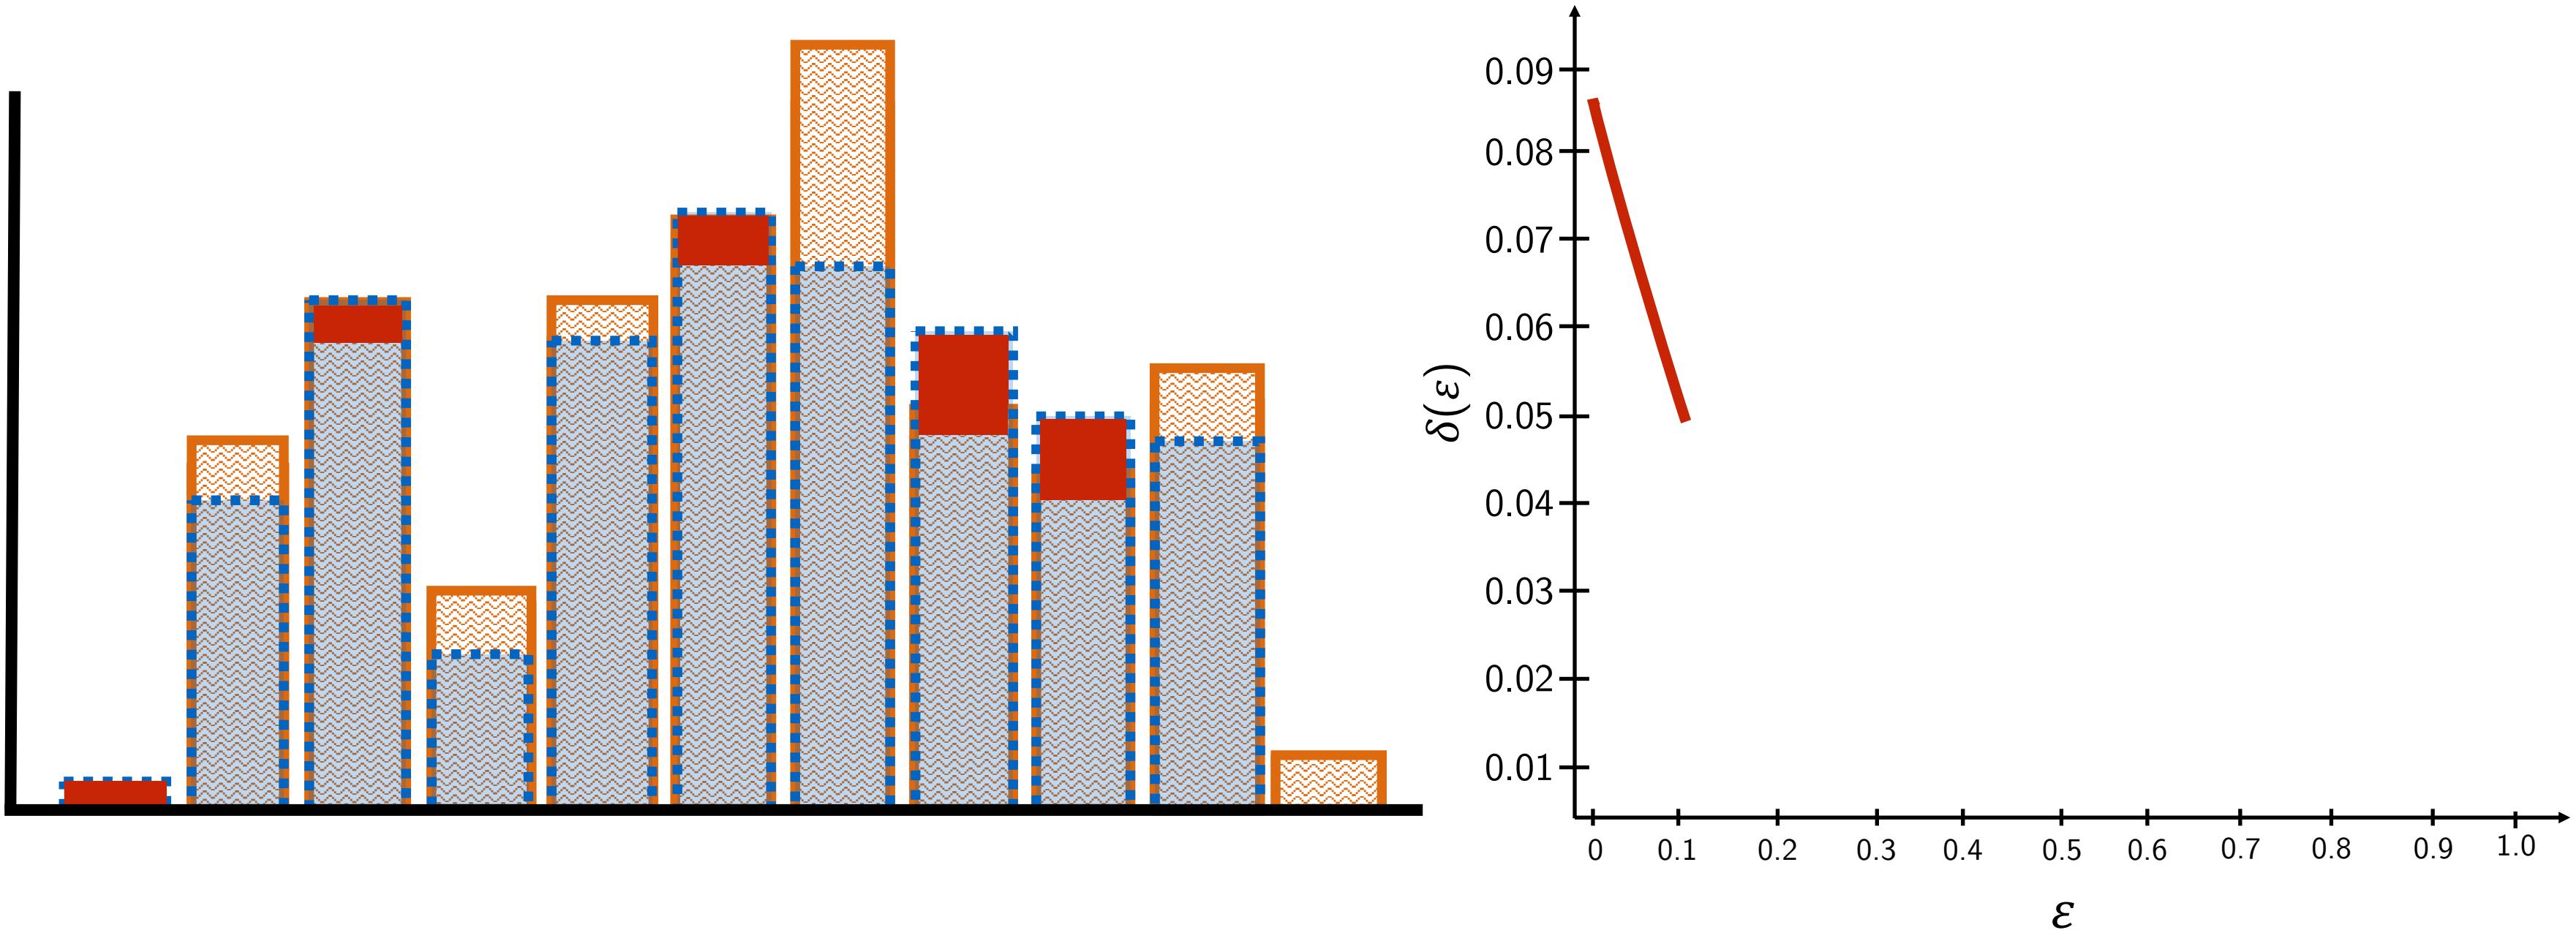
For each ε there is a minimal $\delta(\varepsilon)$

$\varepsilon = 0.1$

$e^\varepsilon = 1.1$

A vs $e^\varepsilon \cdot B$

- Sebastian (A)
- Esfandiar (B)



ε - δ - Graph

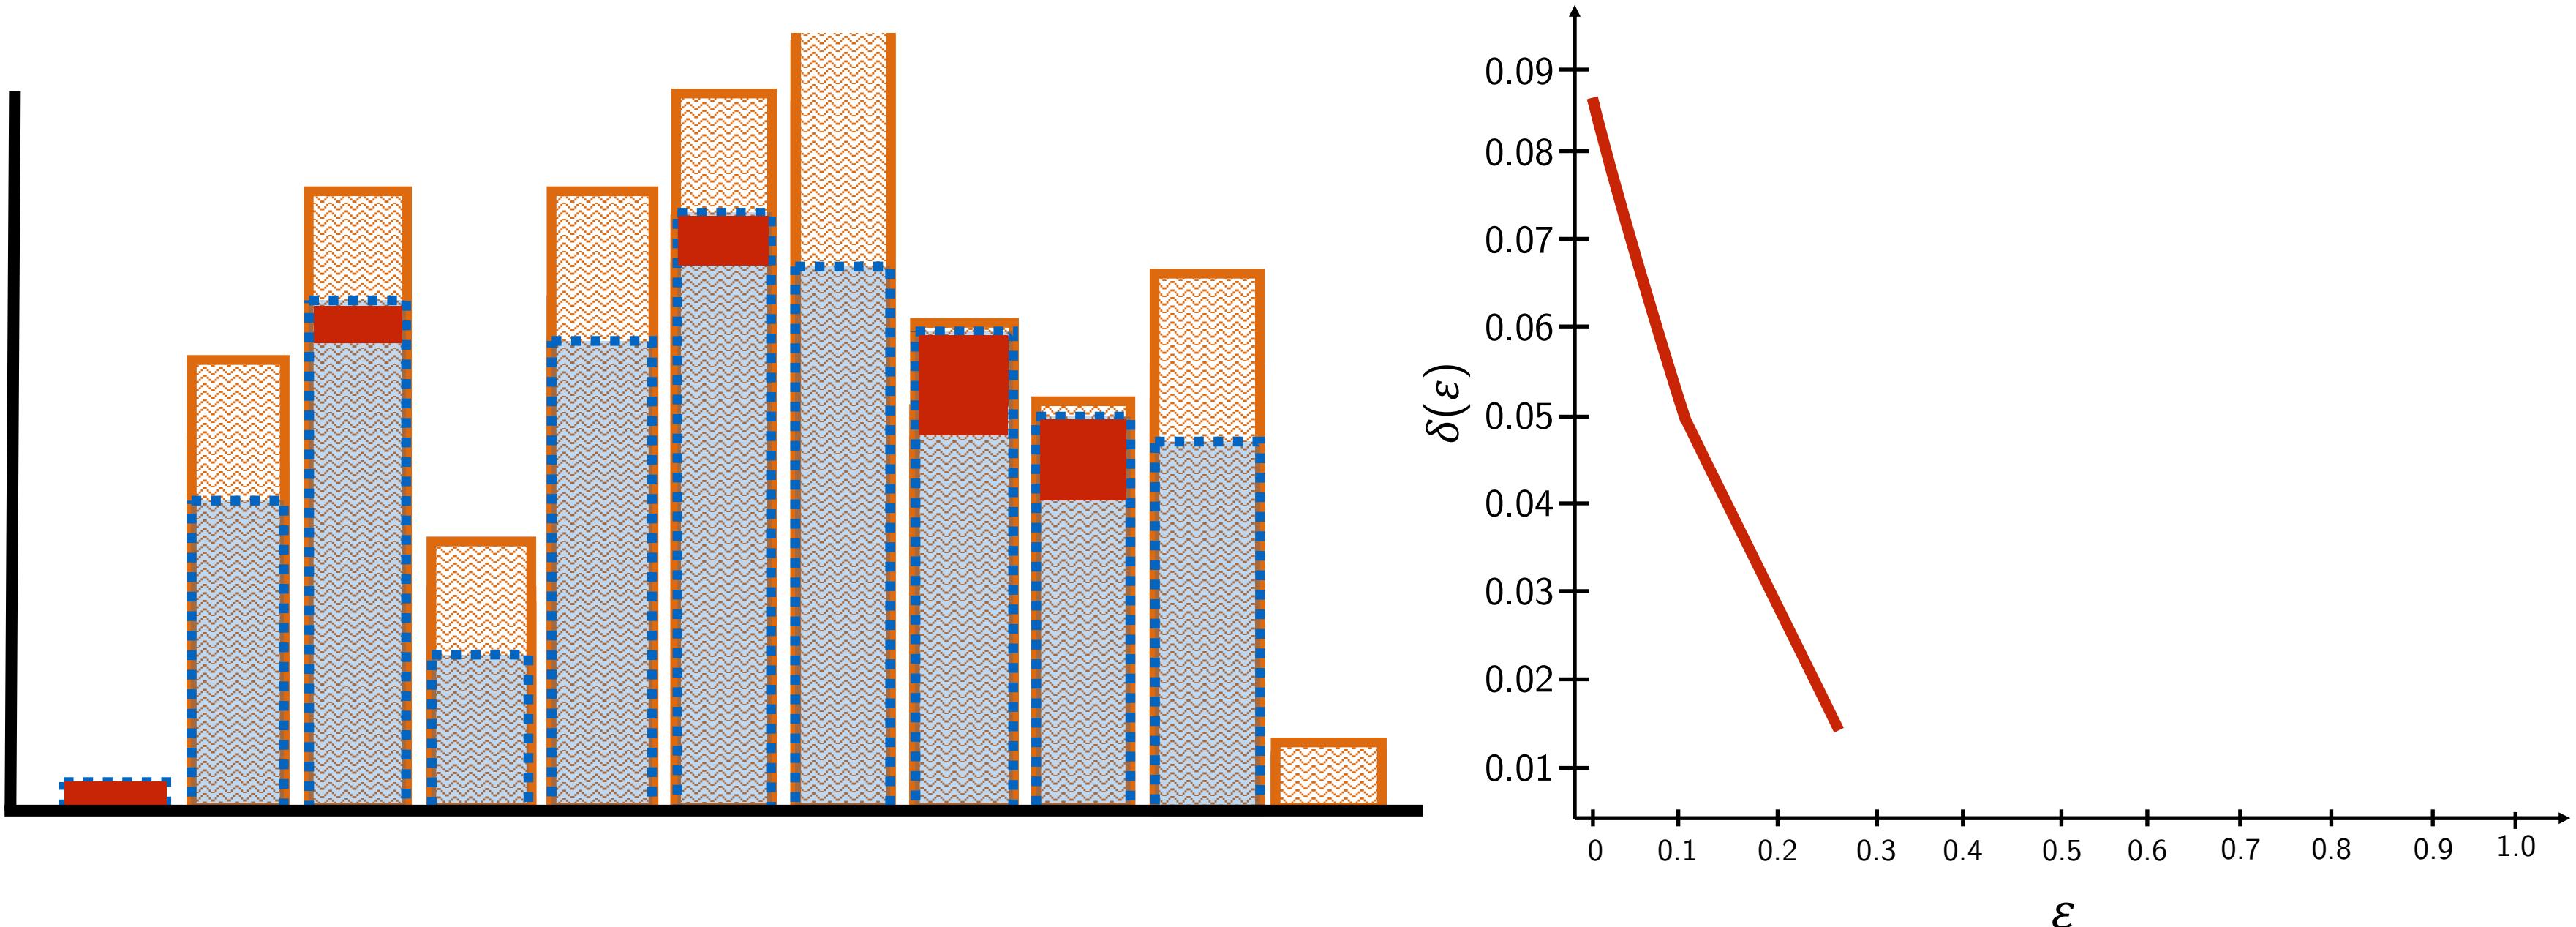
For each ε there is a minimal $\delta(\varepsilon)$

$\varepsilon = 0.3$

$e^\varepsilon = 1.3$

A vs $e^\varepsilon \cdot B$

- Sebastian (A)
- Esfandiar (B)



ε - δ - Graph

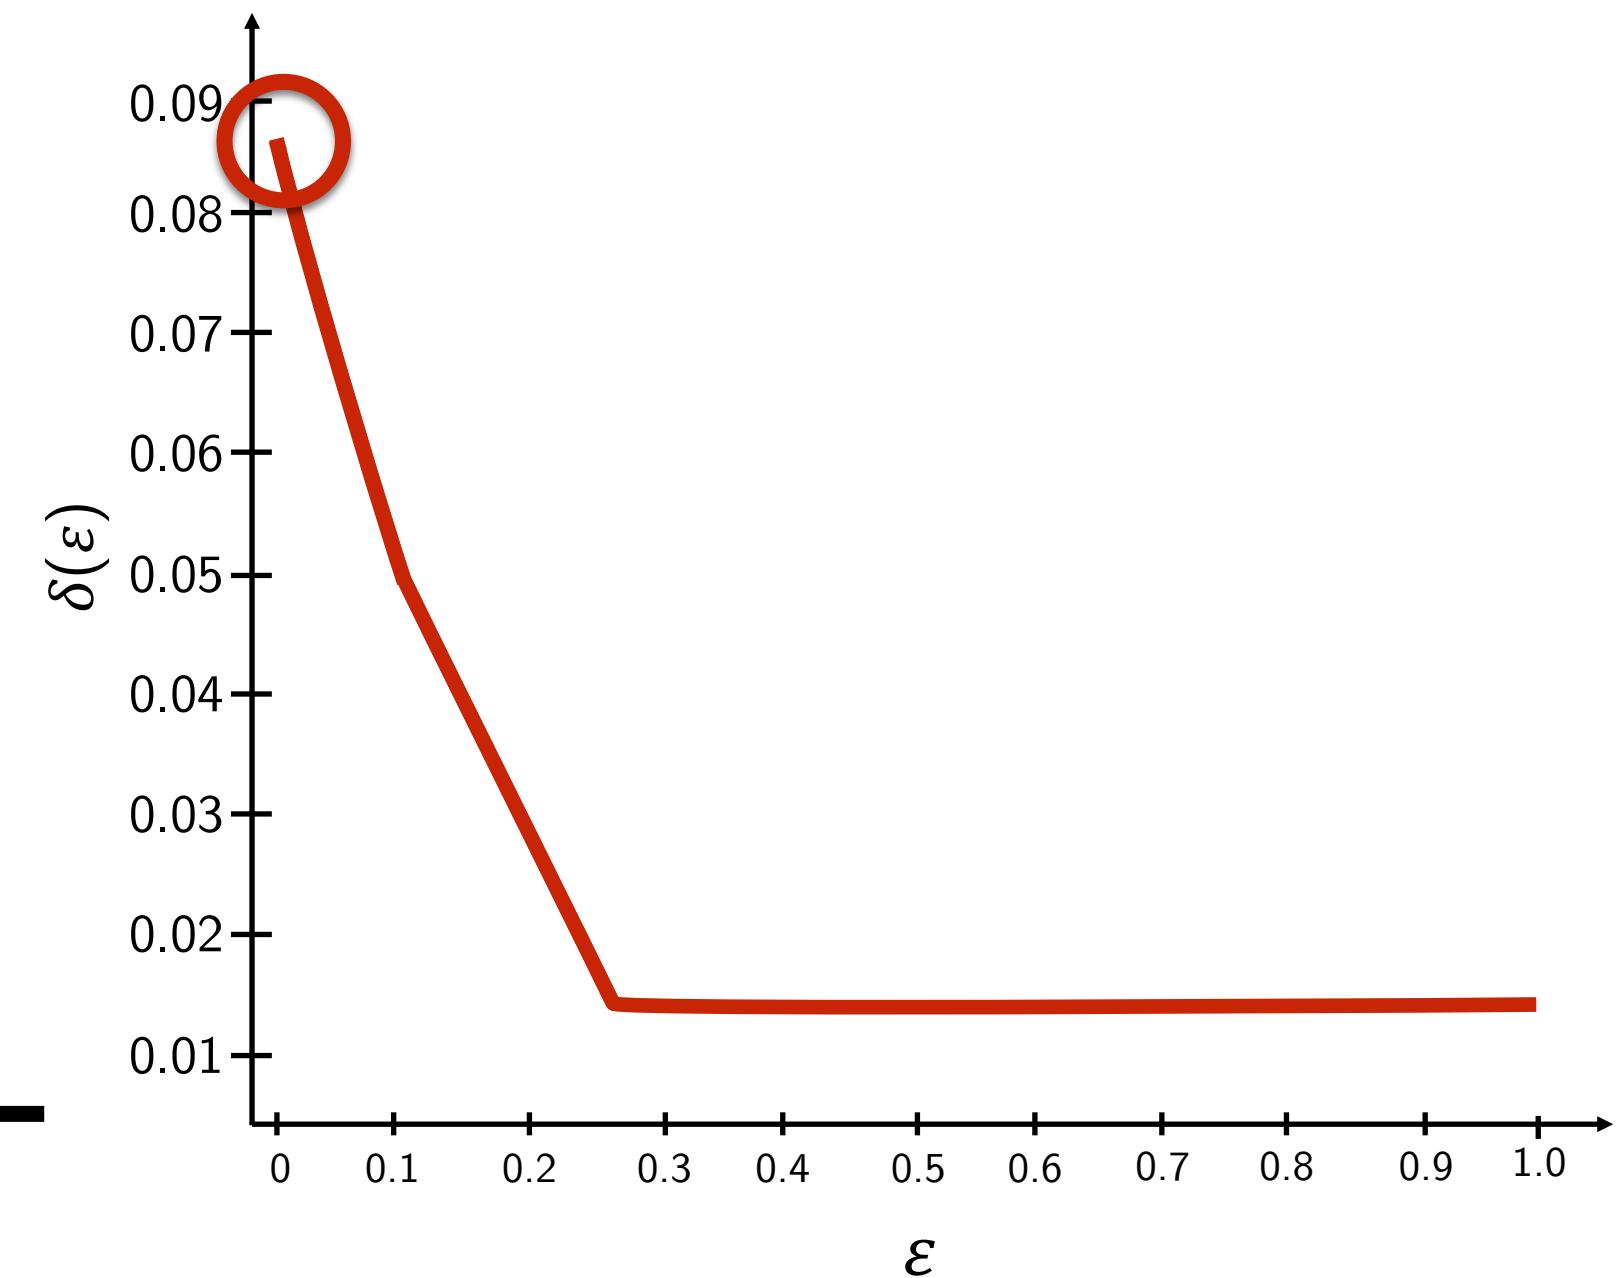
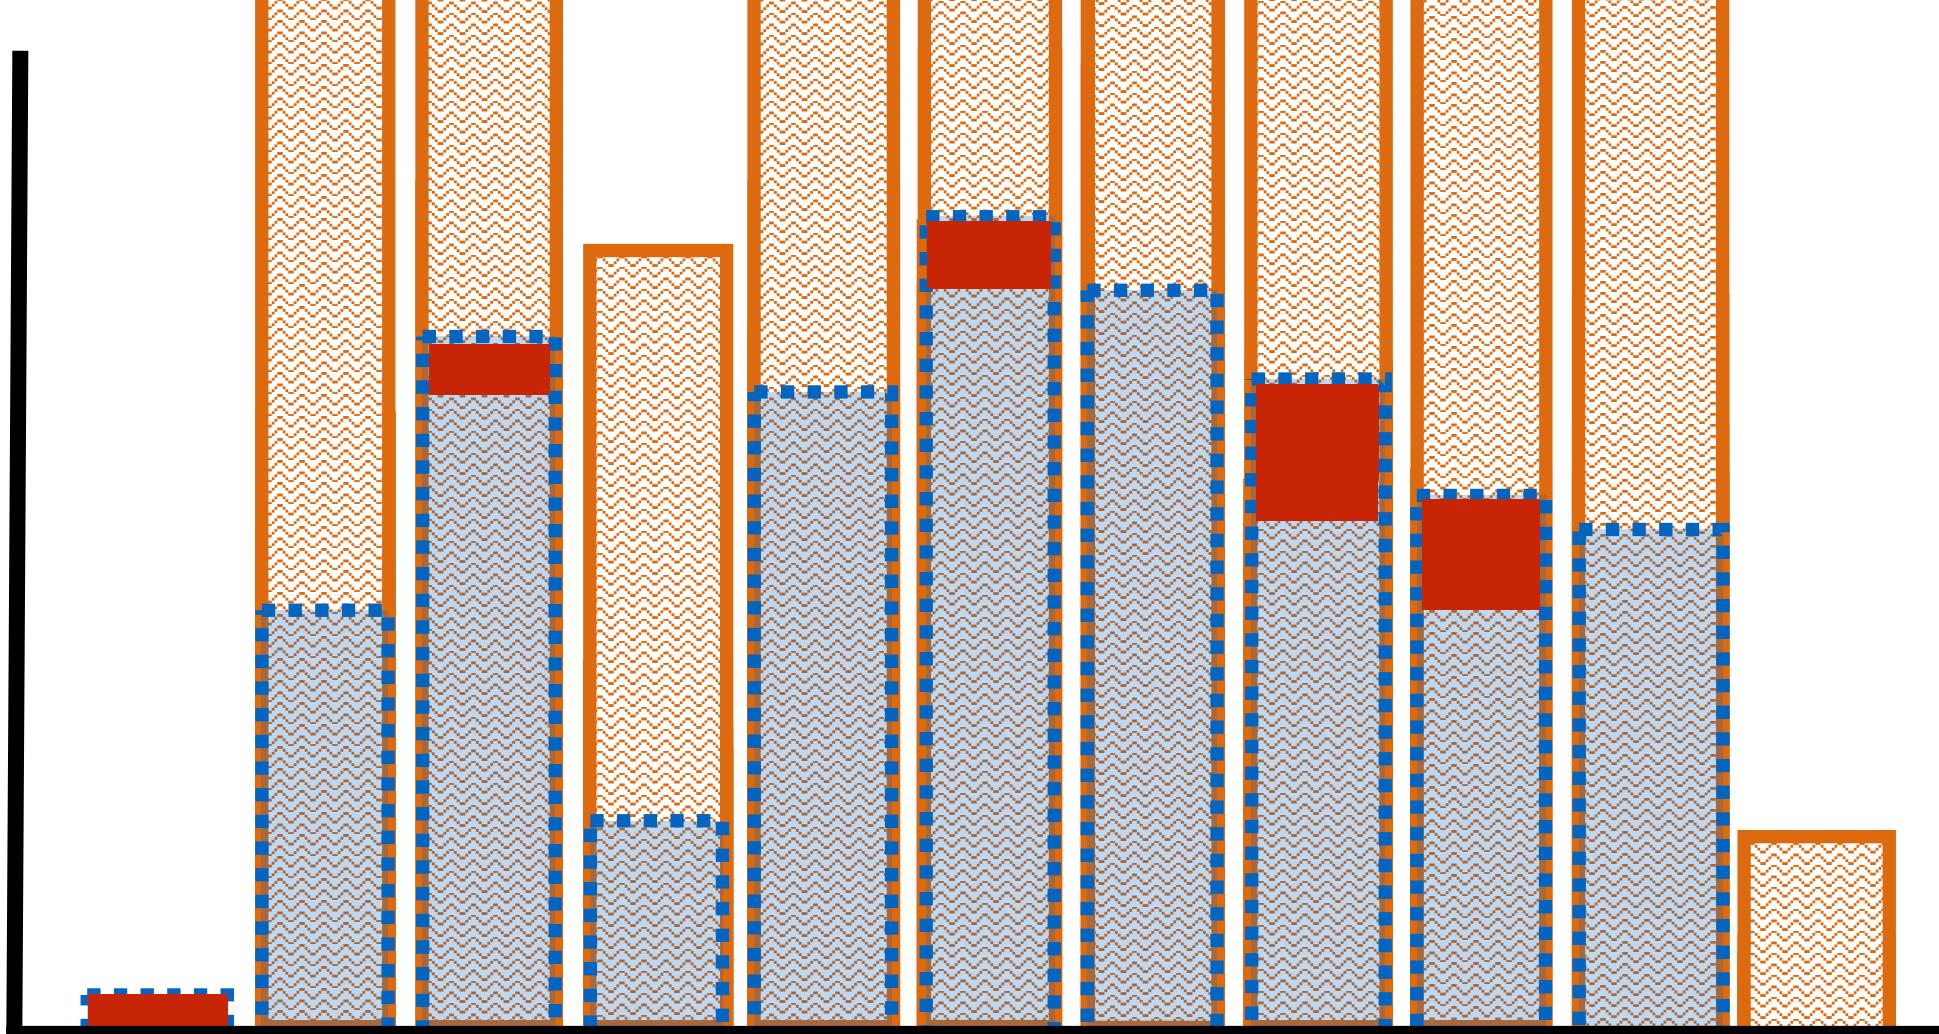
For each ε there is a minimal $\delta(\varepsilon)$

$\varepsilon = 1.0$

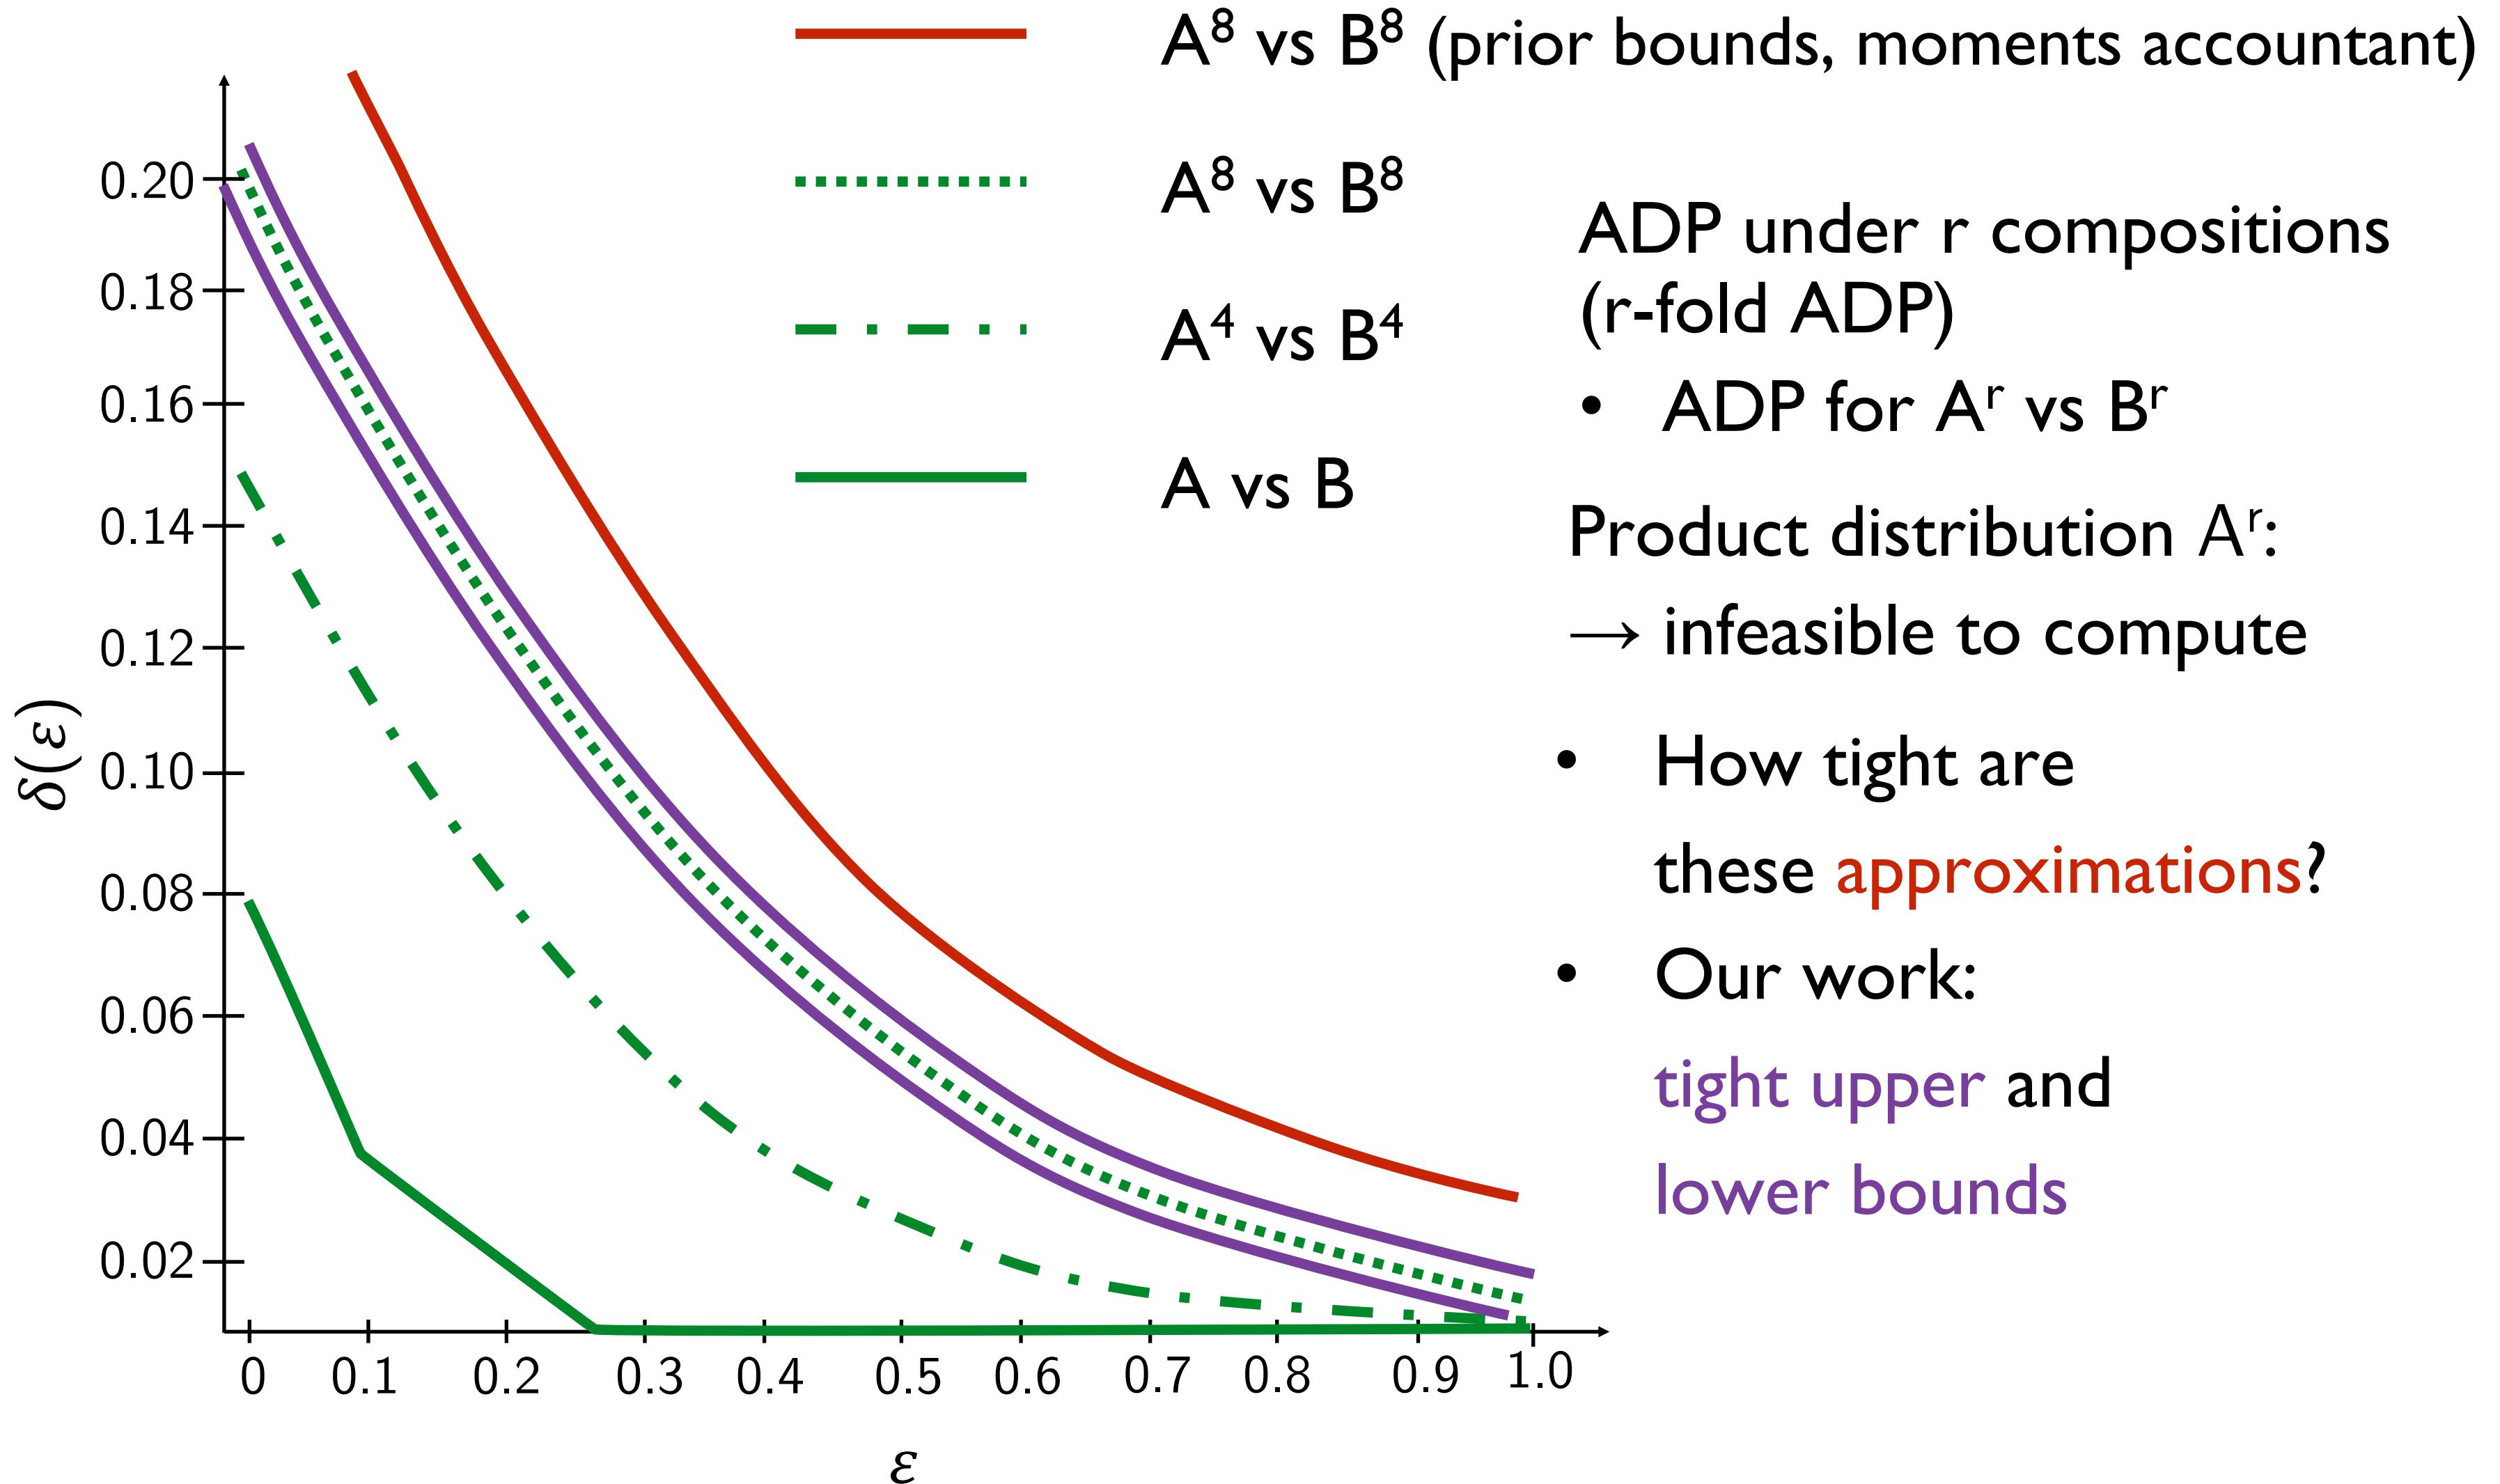
$e^\varepsilon = 2.7$

A vs $e^\varepsilon \cdot B$

- Sebastian (A)
- Esfandiar (B)



R-fold ADP Hard to Compute

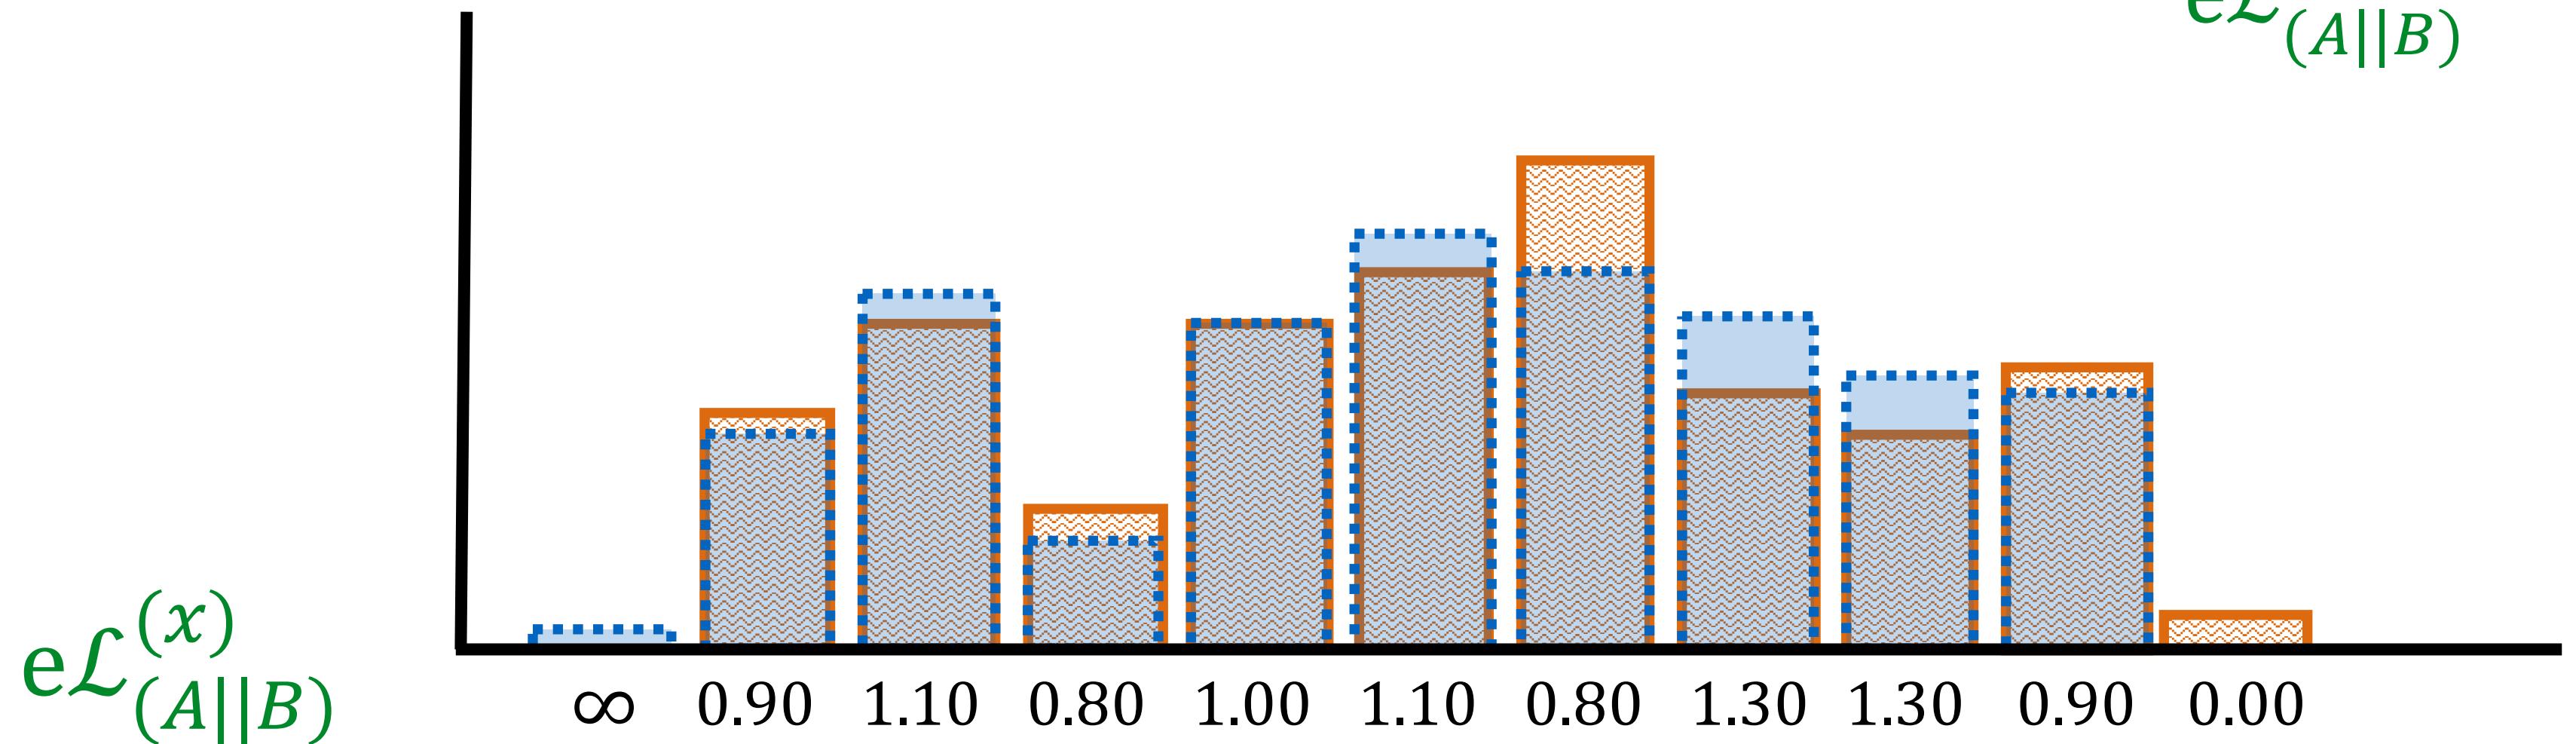


Alternative Approach for Computing r-fold ADP Bounds

$$\begin{aligned}\delta(\varepsilon) &= \sum_x \max(0, \Pr[x \leftarrow A] - e^\varepsilon \Pr[x \leftarrow B]) \\ &= \sum_x \max\left(0, \Pr[x \leftarrow A] \cdot \left(1 - \frac{e^\varepsilon}{e\mathcal{L}_{(A||B)}^{(x)}}\right)\right)\end{aligned}$$

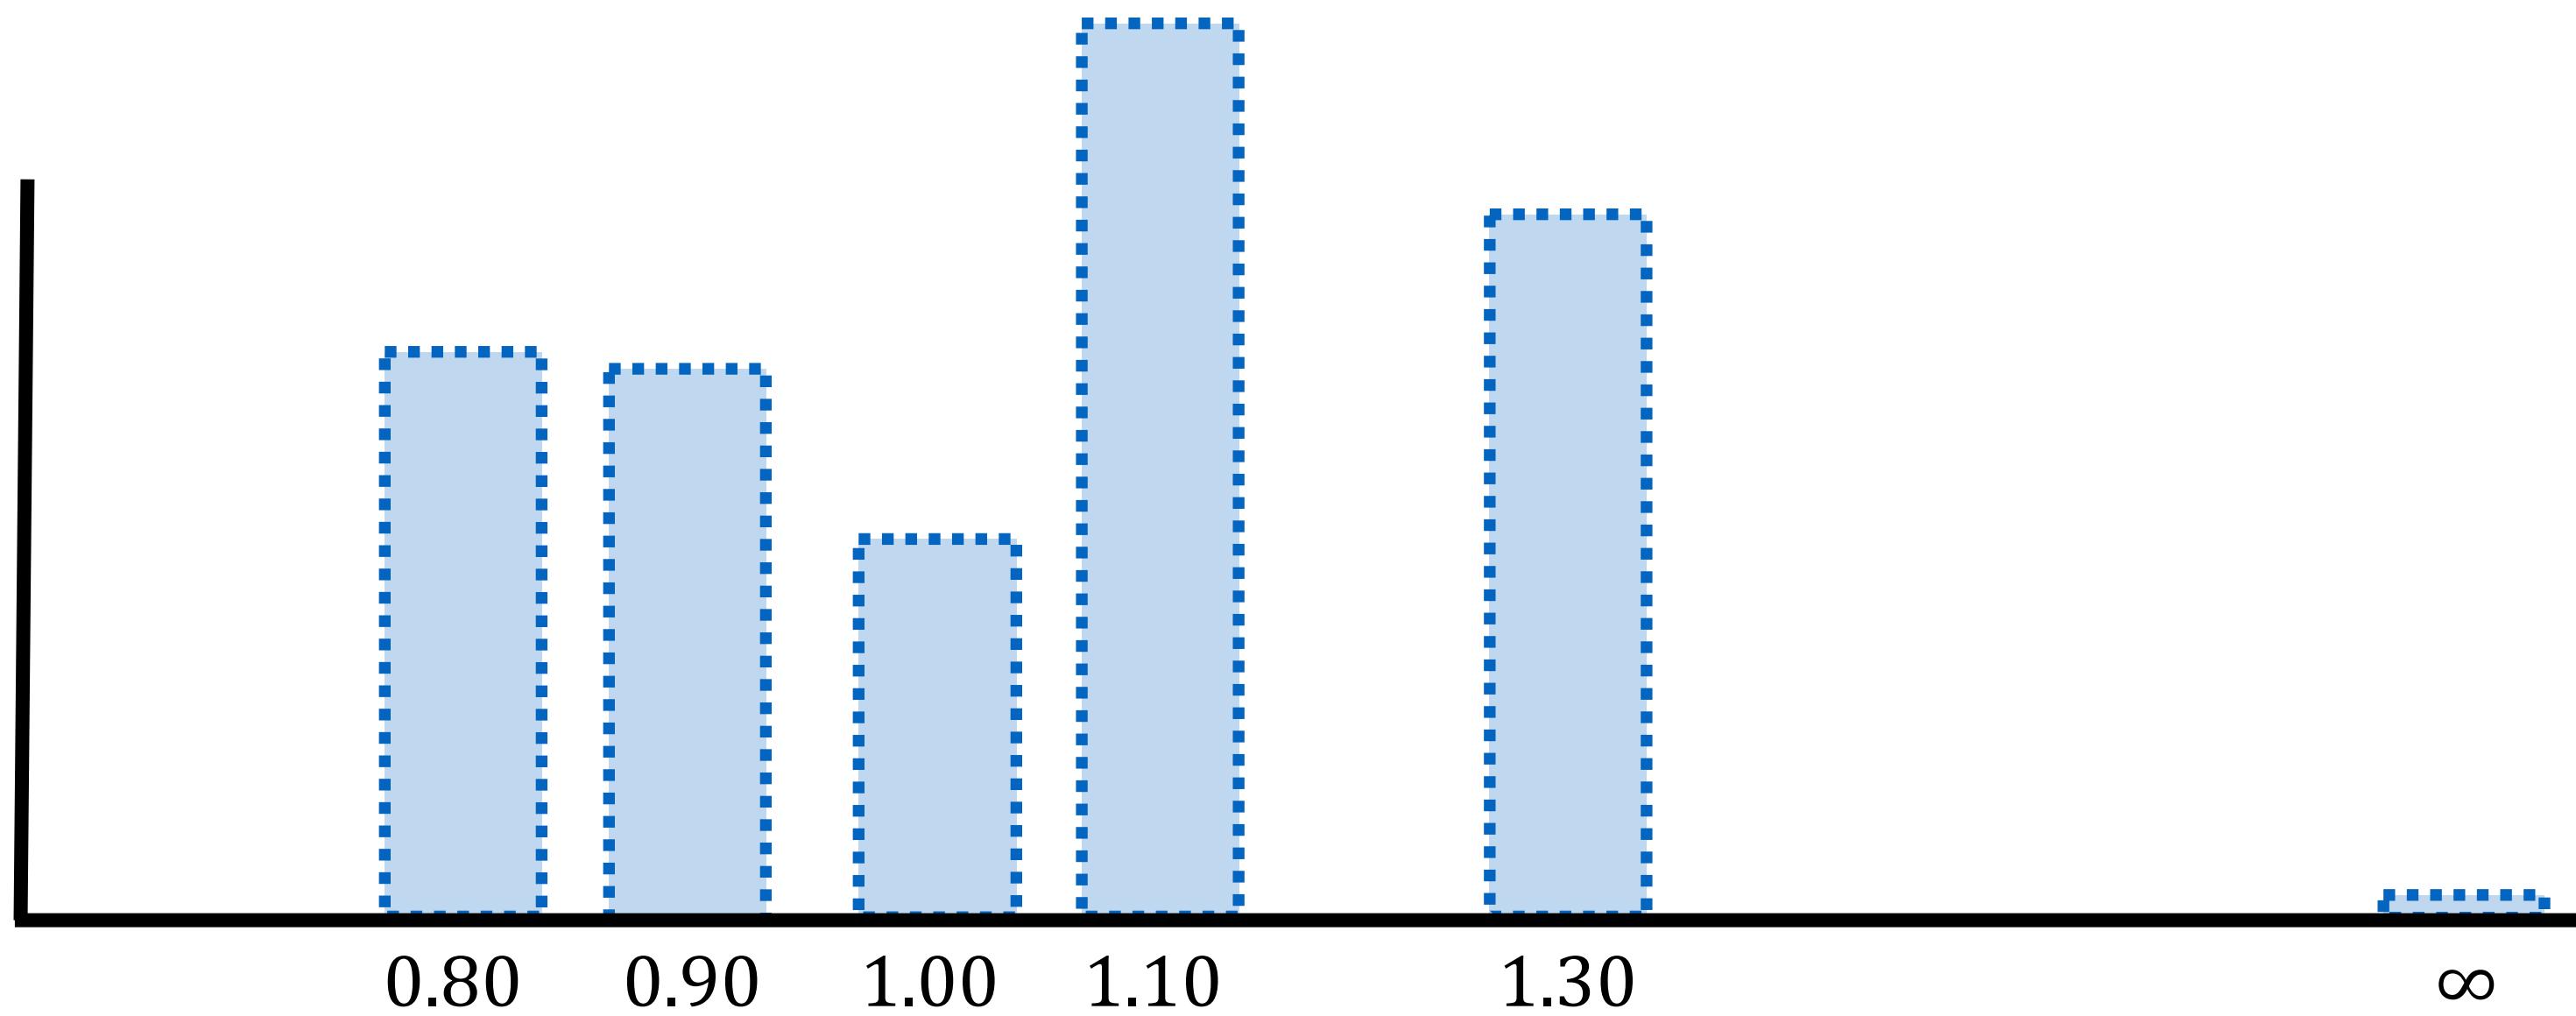
- Sebastian (A)
- Esfandiar (B)

$$\Pr[x \leftarrow B] = \frac{\Pr[x \leftarrow A]}{e\mathcal{L}_{(A||B)}^{(x)}}$$



$$\text{Privacy loss } e\mathcal{L}_{(A||B)}^{(x)} = \frac{\Pr[x \leftarrow A]}{\Pr[x \leftarrow B]}$$

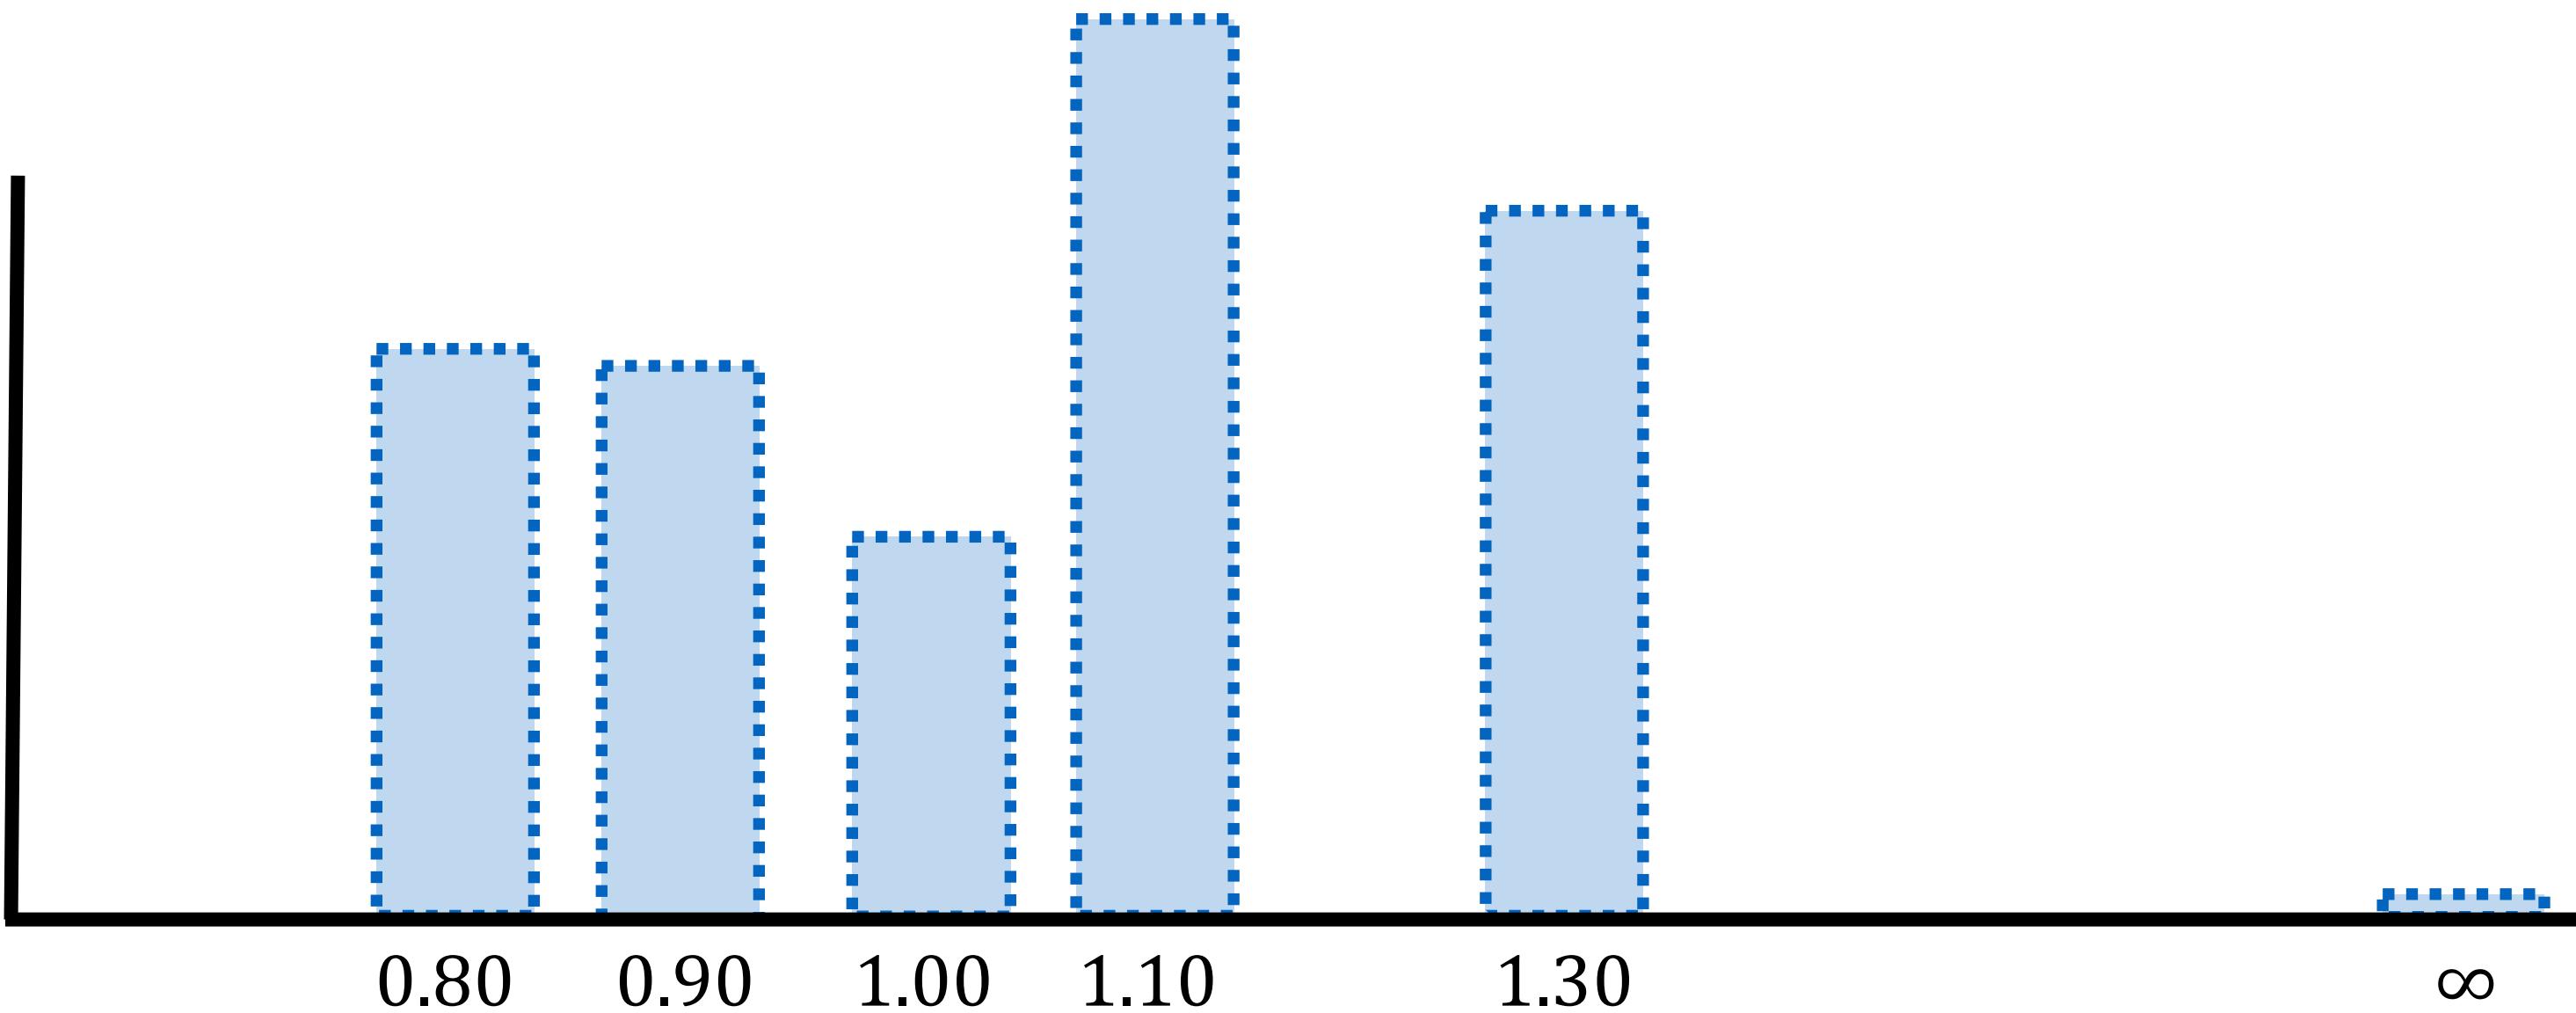
Summarizing and ordering by the privacy loss



Composing the privacy loss

$$e\mathcal{L}_{(A||B)}^{(x)} \quad e\mathcal{L}_{(A||B)}^{(y)}$$

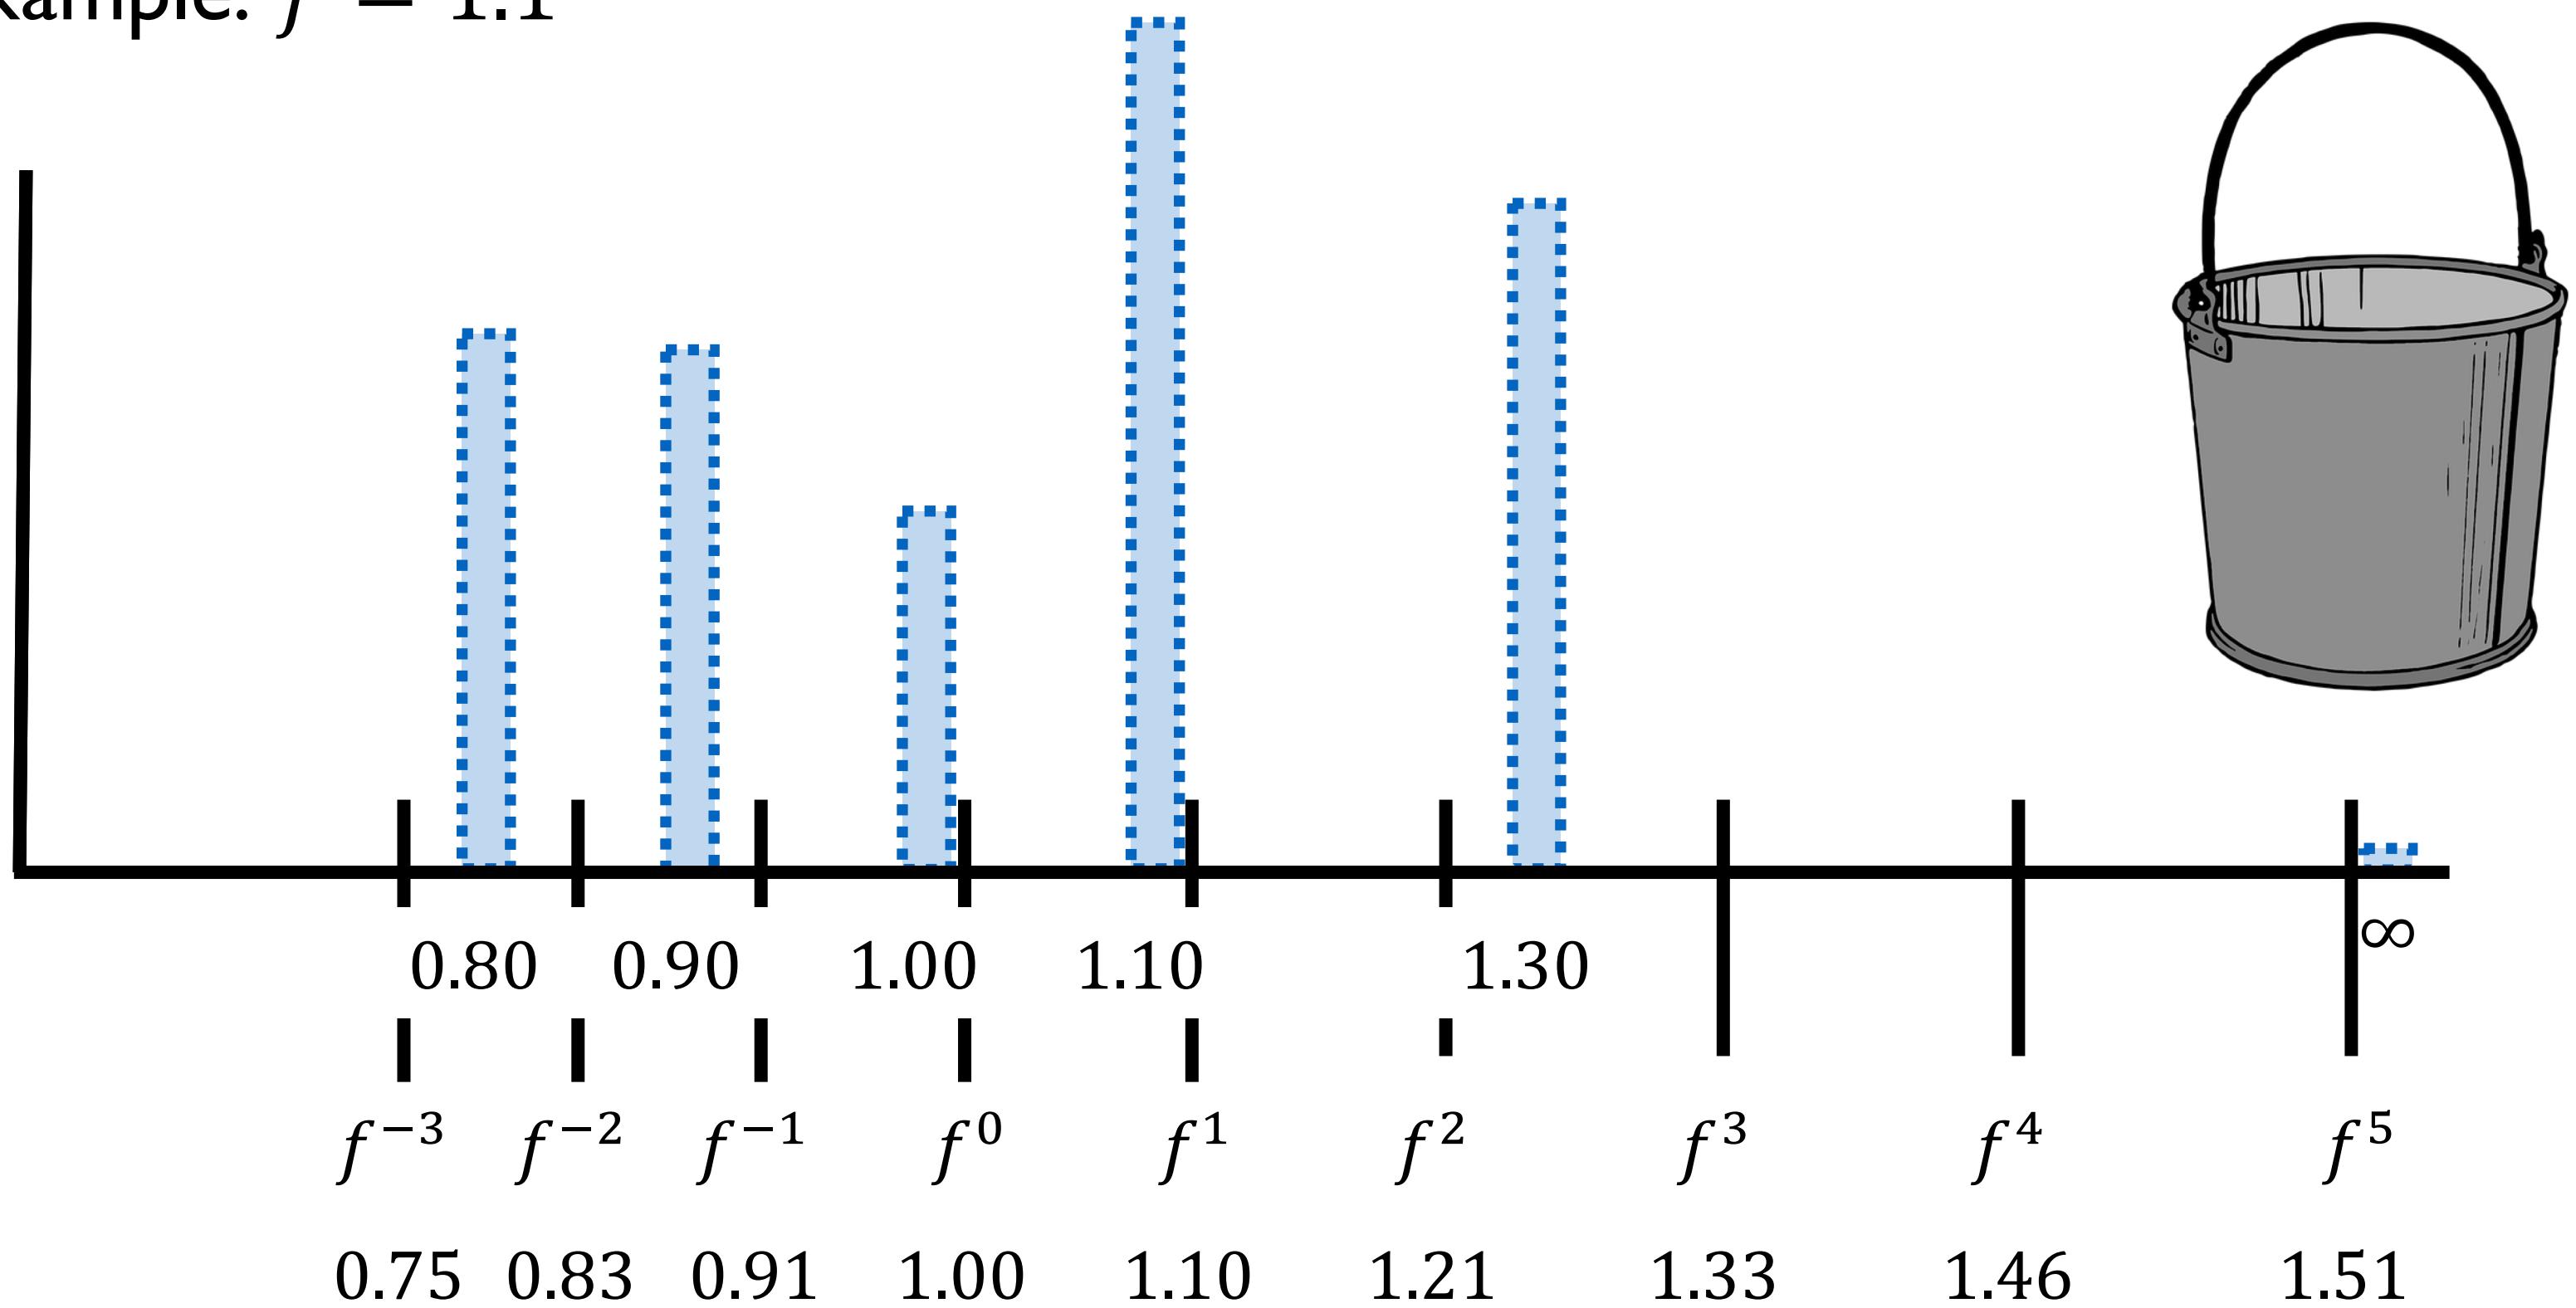
$$e\mathcal{L}_{(A^2||B^2)}^{(x,y)} = \frac{\Pr[(x,y) \leftarrow A^2]}{\Pr[(x,y) \leftarrow B^2]} = \frac{\Pr[x \leftarrow A]}{\Pr[x \leftarrow B]} \cdot \frac{\Pr[y \leftarrow A]}{\Pr[y \leftarrow B]}$$



Privacy Buckets

Buckets for making composition easier:

- Approximate privacy loss to f^i for different i (always round up)
- Example: $f = 1.1$



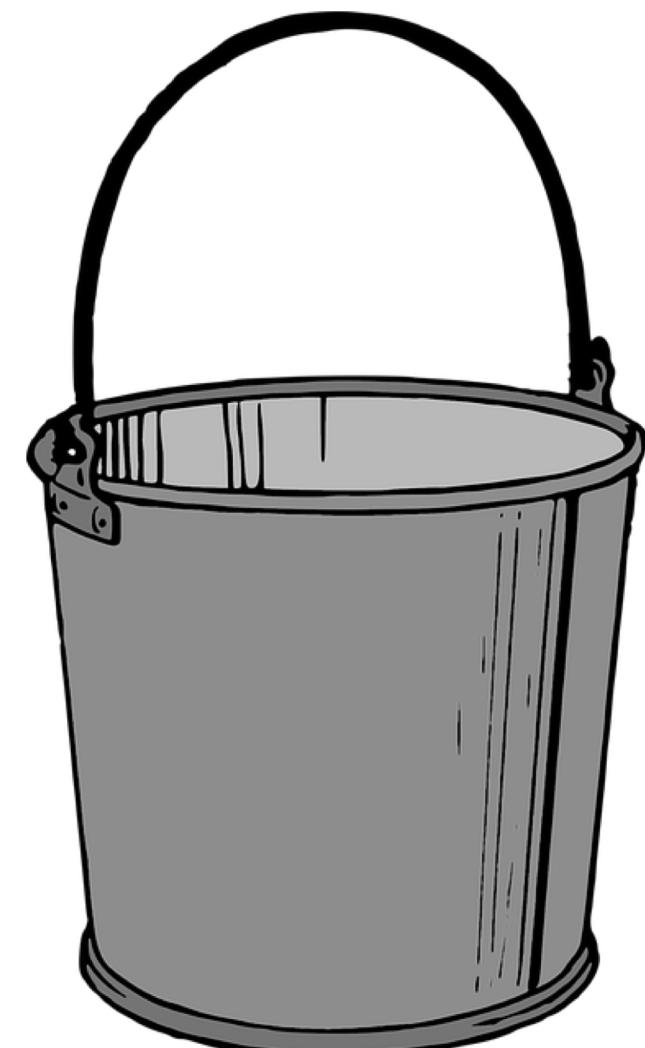
Precision of our privacy buckets approximation

What do we actually do and how large is the error?

$$e\mathcal{L}_{(A||B)}^{(x)} = \frac{\Pr[\text{Seb. watches } x]}{\Pr[\text{Esf. watches } x]} = \frac{\boxed{}}{\boxed{}}$$

$f^{i-1} \quad f^i \quad f^{i+1}$

We later reconstruct $\boxed{}$ as $\frac{\Pr[\boxed{}]}{f^i}$ instead of $\frac{\Pr[\boxed{}]}{e\mathcal{L}_{(A||B)}^{(x)}}$

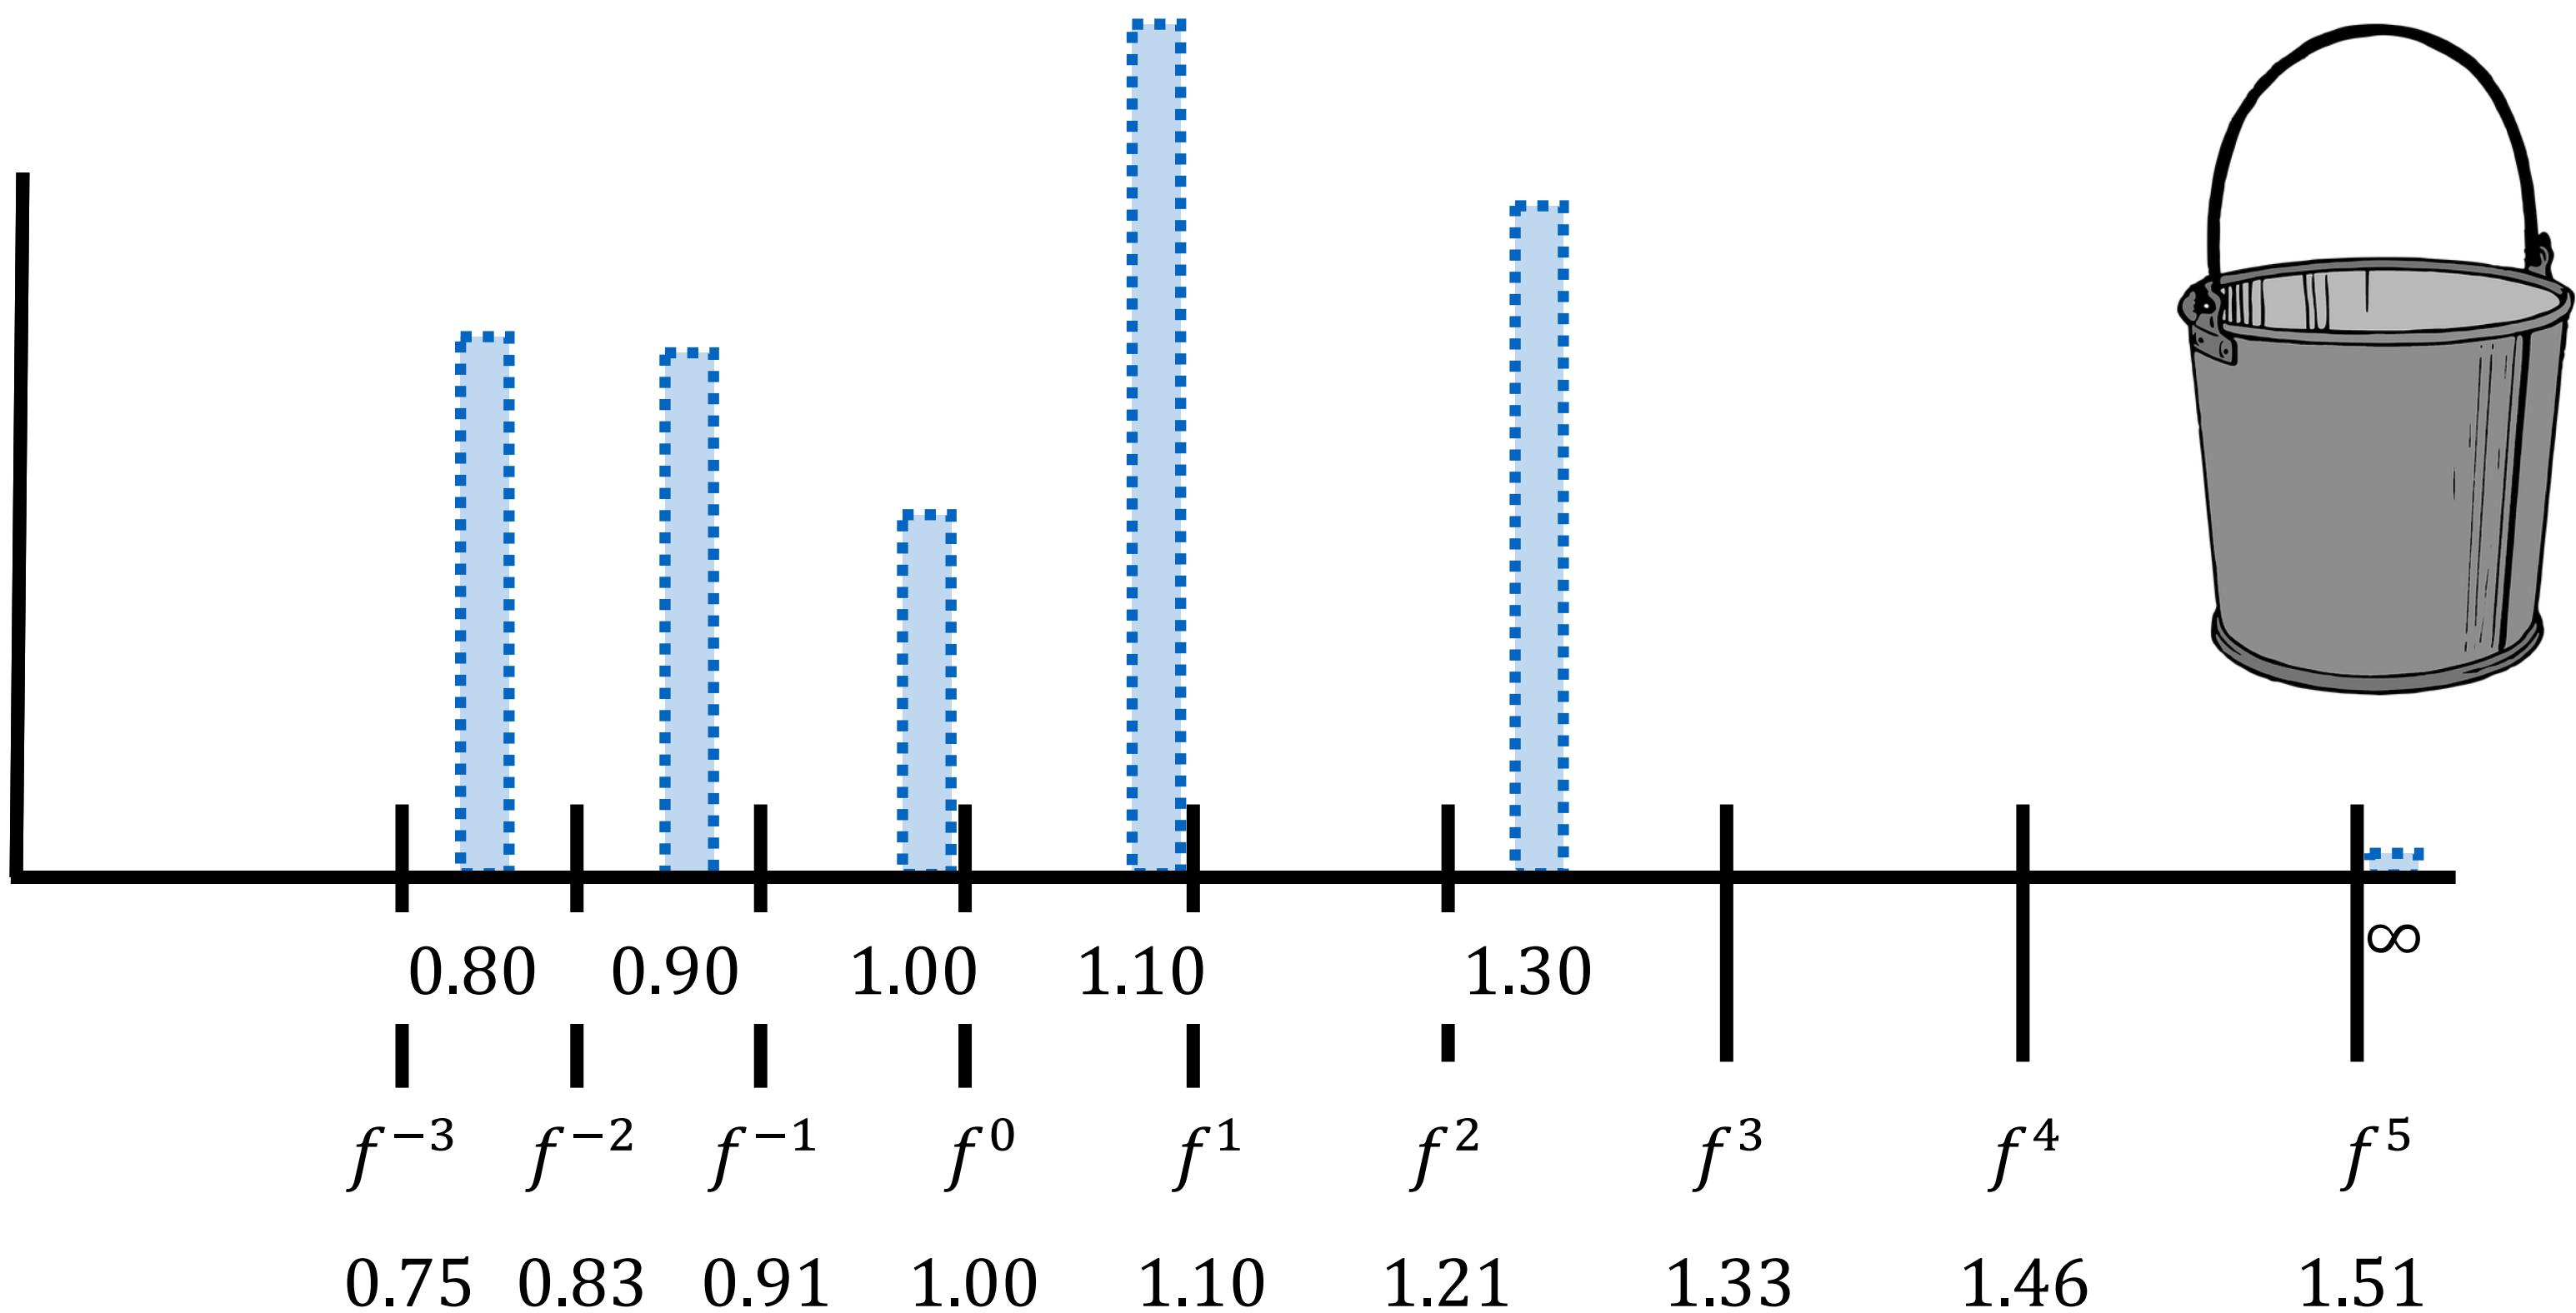


Privacy Buckets: composition

Buckets with $f = 1.1$

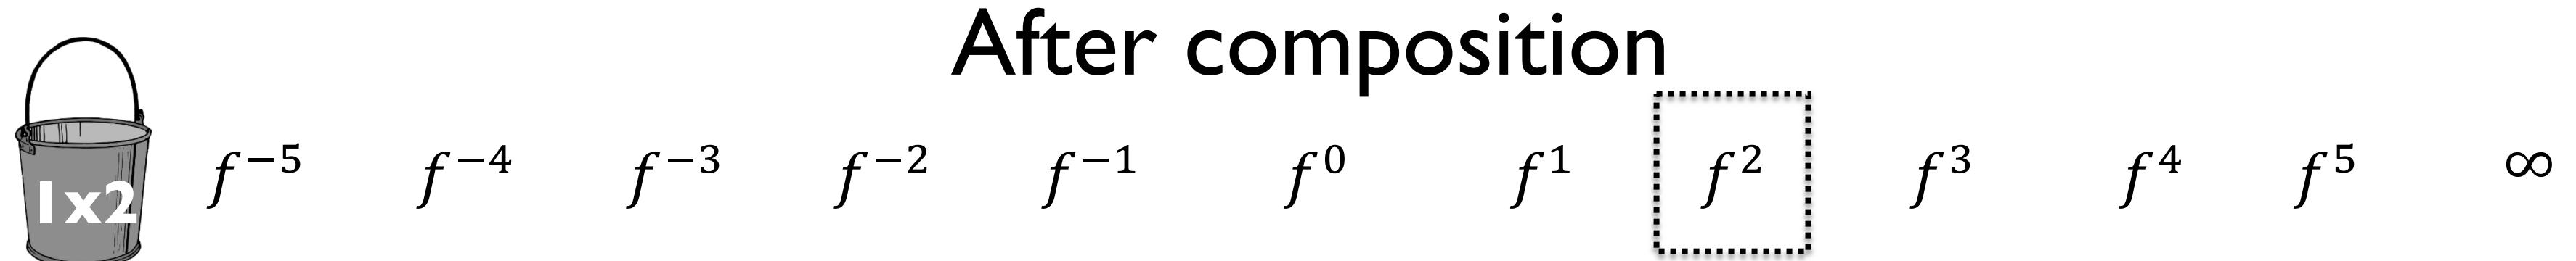
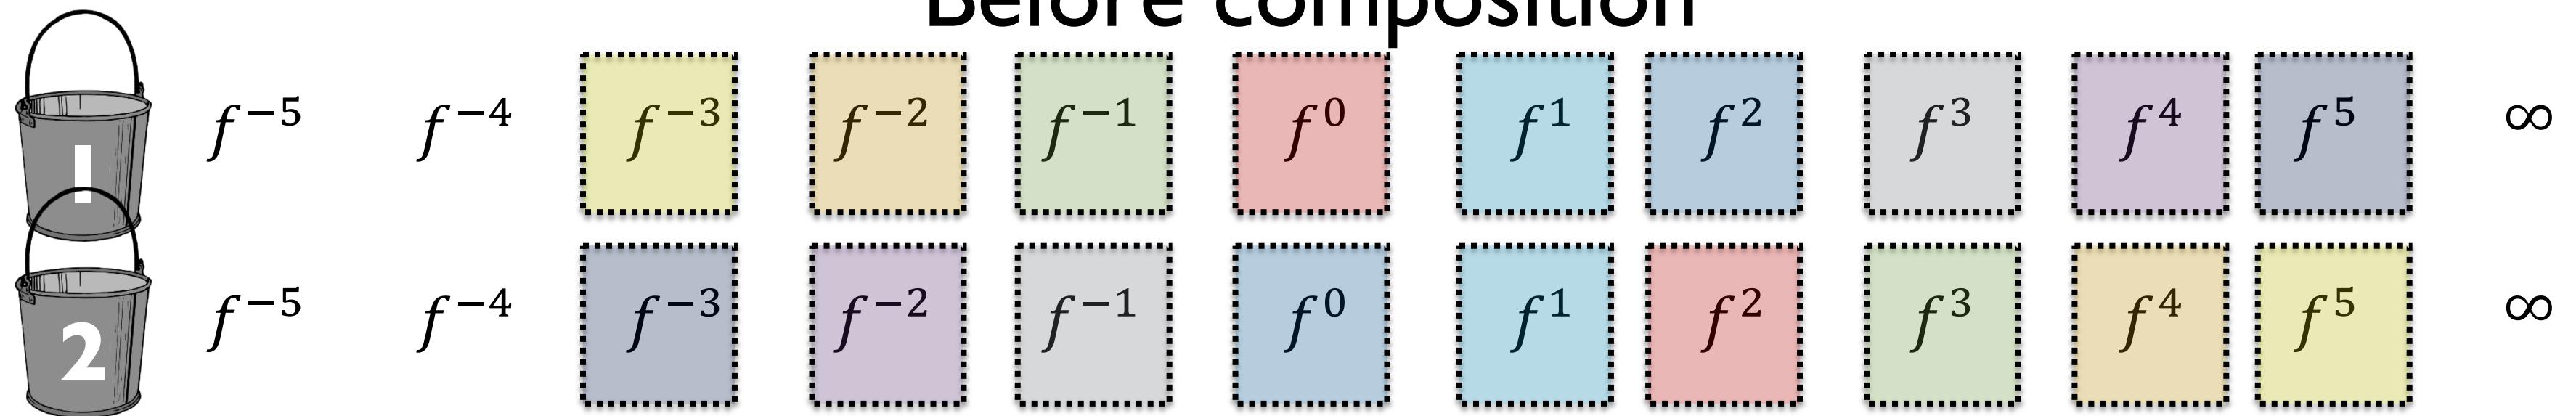
$$f^{j+i} \quad f^j \quad f^i$$

$$e\mathcal{L}_{(A^2||B^2)}^{(x,y)} = \frac{\Pr[(x,y) \leftarrow A^2]}{\Pr[(x,y) \leftarrow B^2]} = \frac{\Pr[x \leftarrow A]}{\Pr[x \leftarrow B]} \cdot \frac{\Pr[y \leftarrow A]}{\Pr[y \leftarrow B]}$$



Privacy buckets composition

Example: we compute f^2 for buckets from f^{-5} to f^5 (and ∞)

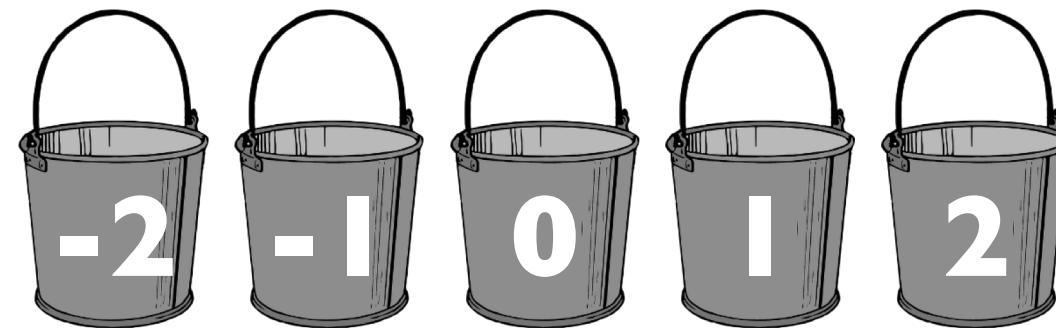


$$I \times 2(i) = \sum_{j,k \text{ s.t. } j+k=i} I(j) \cdot 2(k)$$

Feasibility

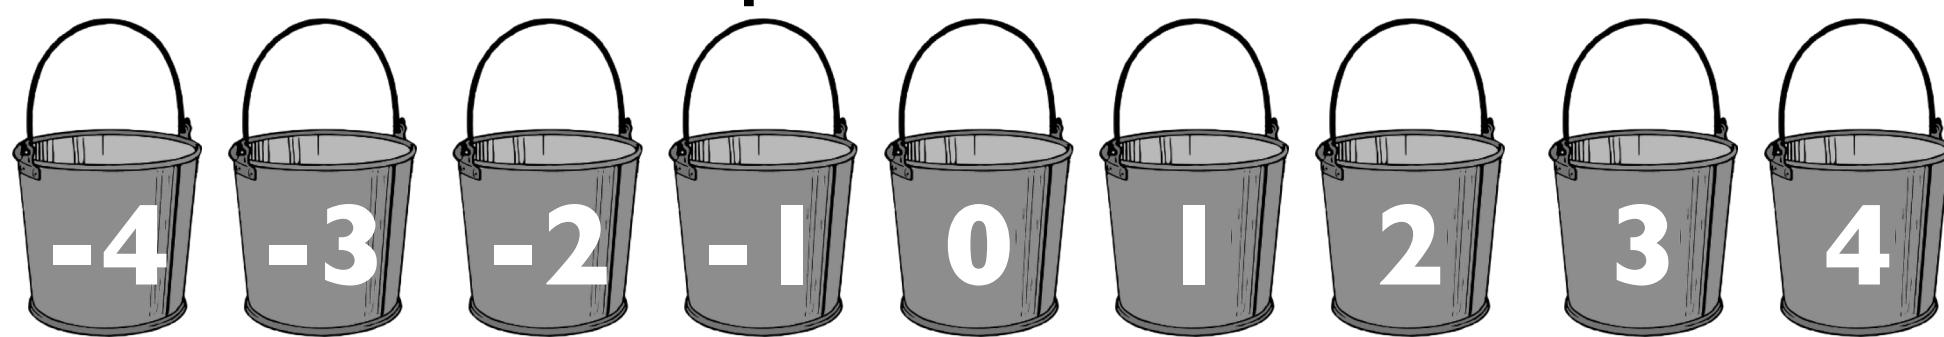
How does composition affect the number of required buckets?

Starting with 5 buckets

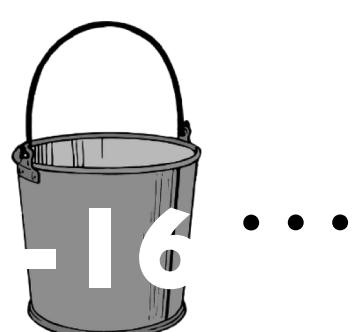


2 observations

1 self-composition \rightarrow 9 buckets



Collect all $\leq -n$



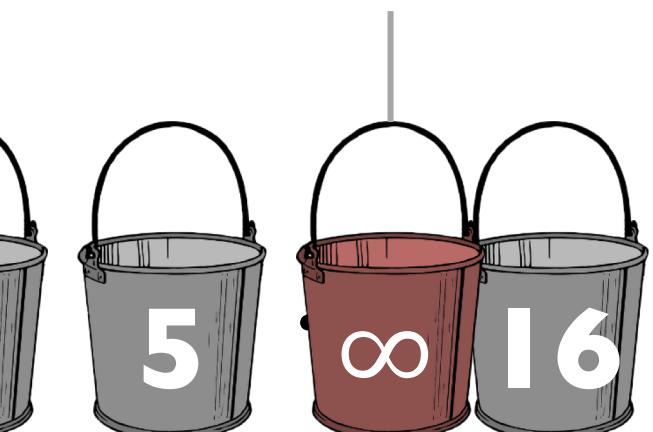
4 observations

2 self-compositions

limit to $2n+2$ buckets!



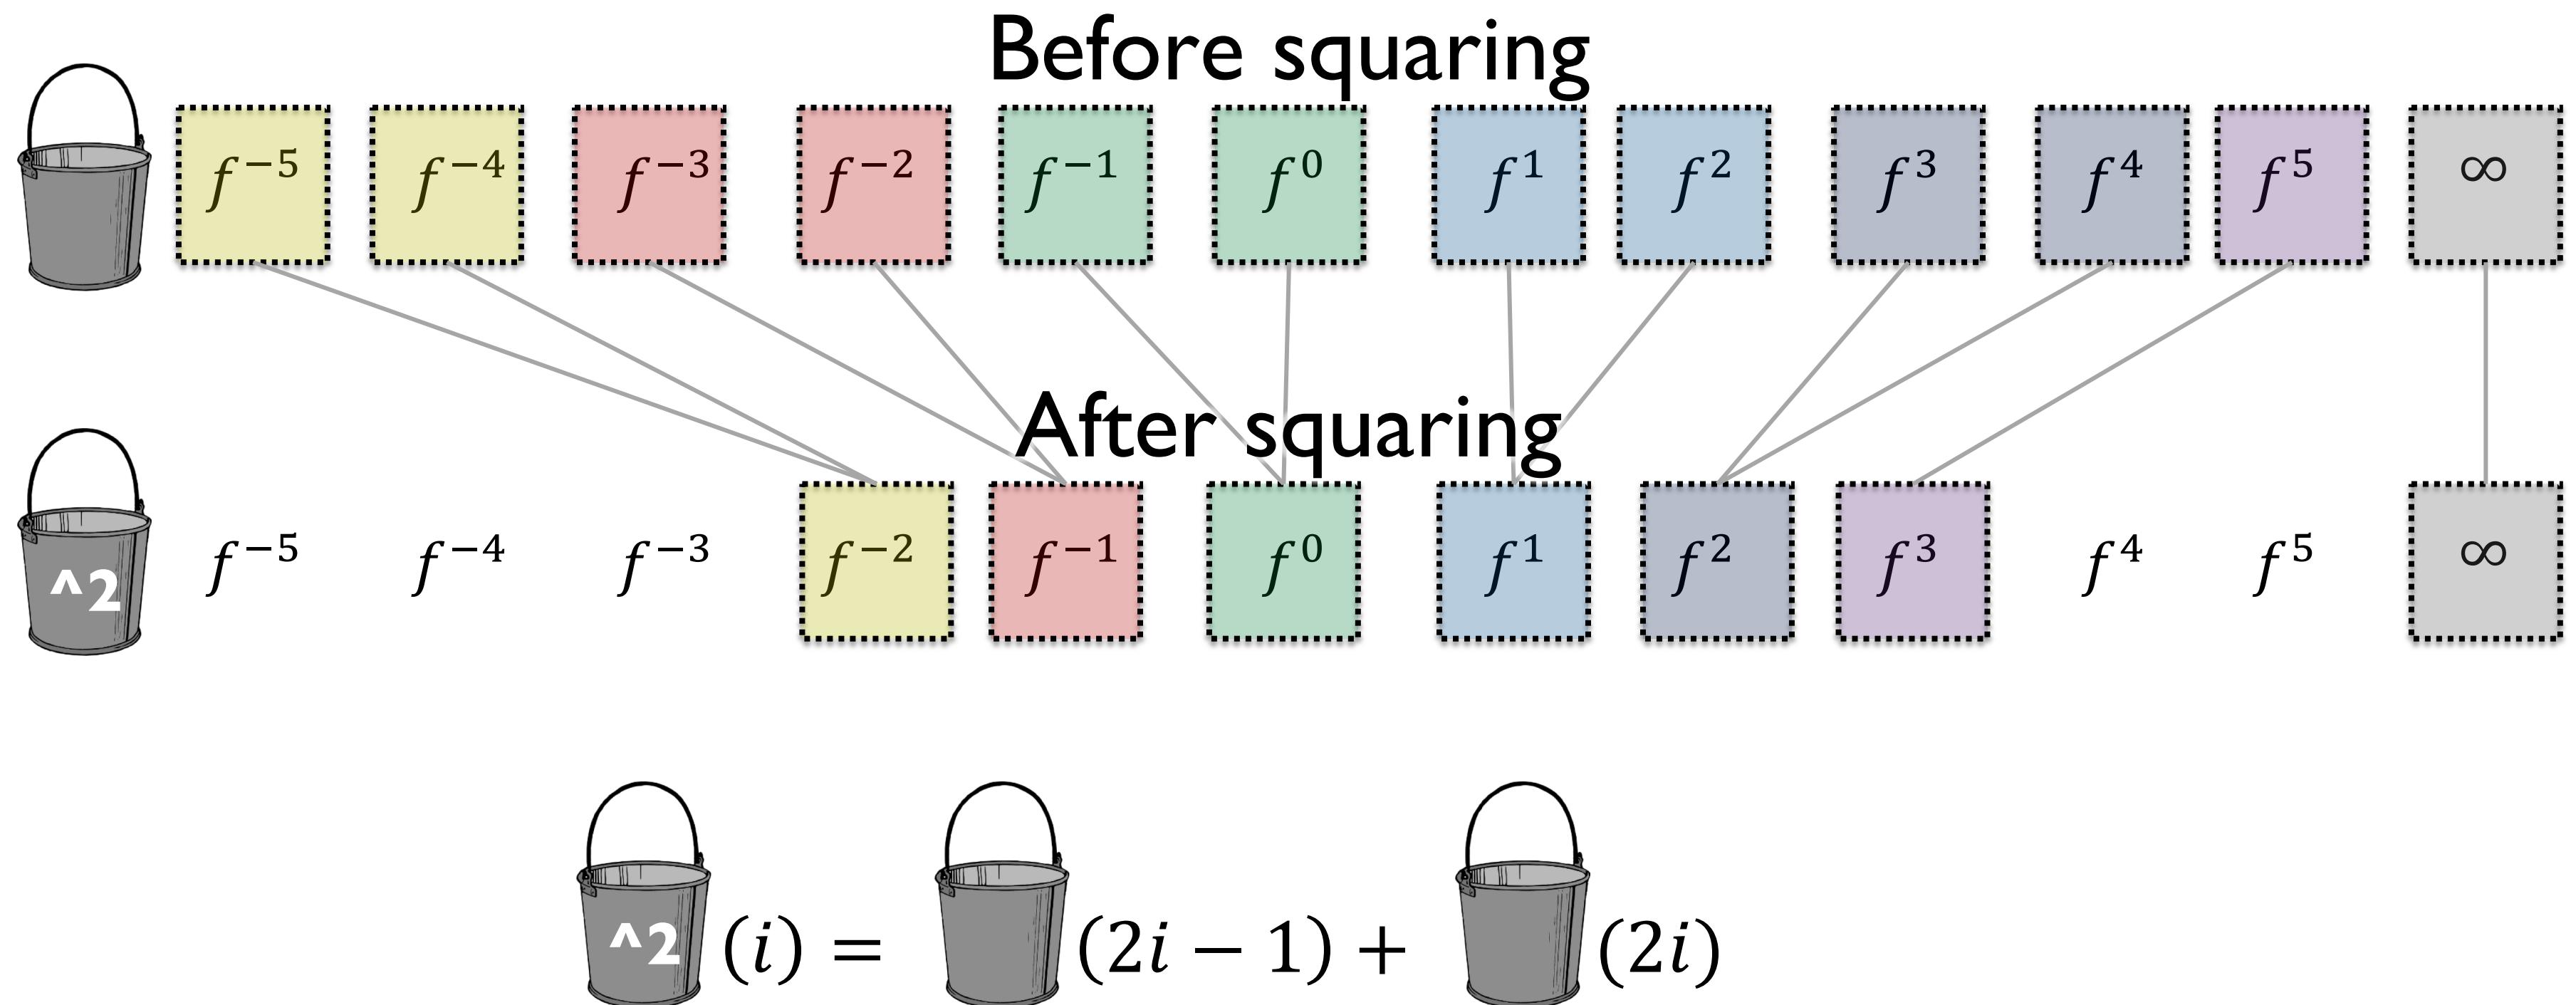
Collect all $> n$



Privacy buckets squaring

We “square” the buckets, i.e., we square f to f^2 .

We then rearrange the buckets.



Implementation and Performance

Unoptimized Python (NumPy) implementation
(405 LoC) accessible under:

<https://github.com/sommerda/privacybuckets>

- Reimplemented by David Sommer (ETH)

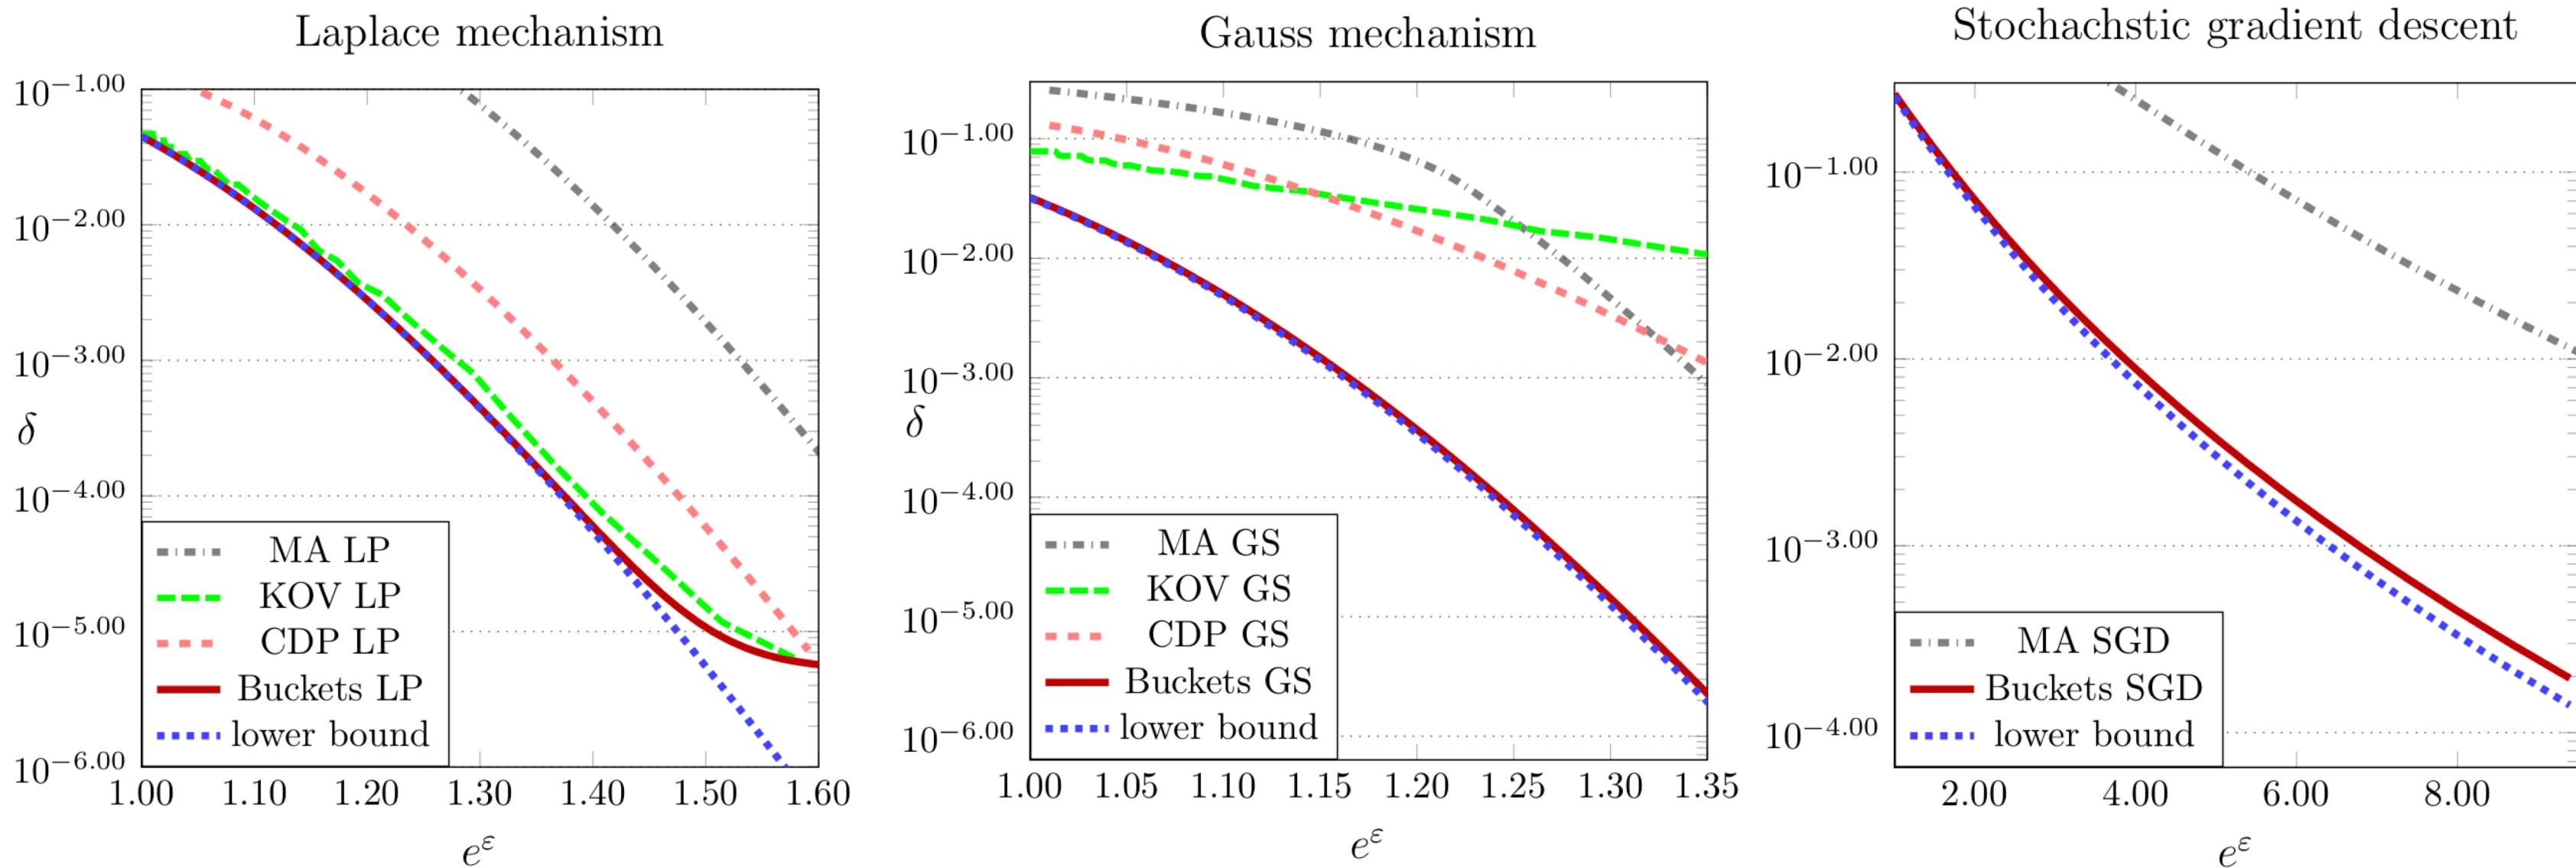
Computation time dominated by convolution:

- FFT-based convolution $O(n \log n)$
- Implementation uses normal convolution $O(n^2)$ for numerical stability
- in total ~ 13 seconds per composition operation
(with 100 000 buckets)
- Repeated squaring: 2^r -fold ADP bounds requires r compositions

How tight are these
bounds?

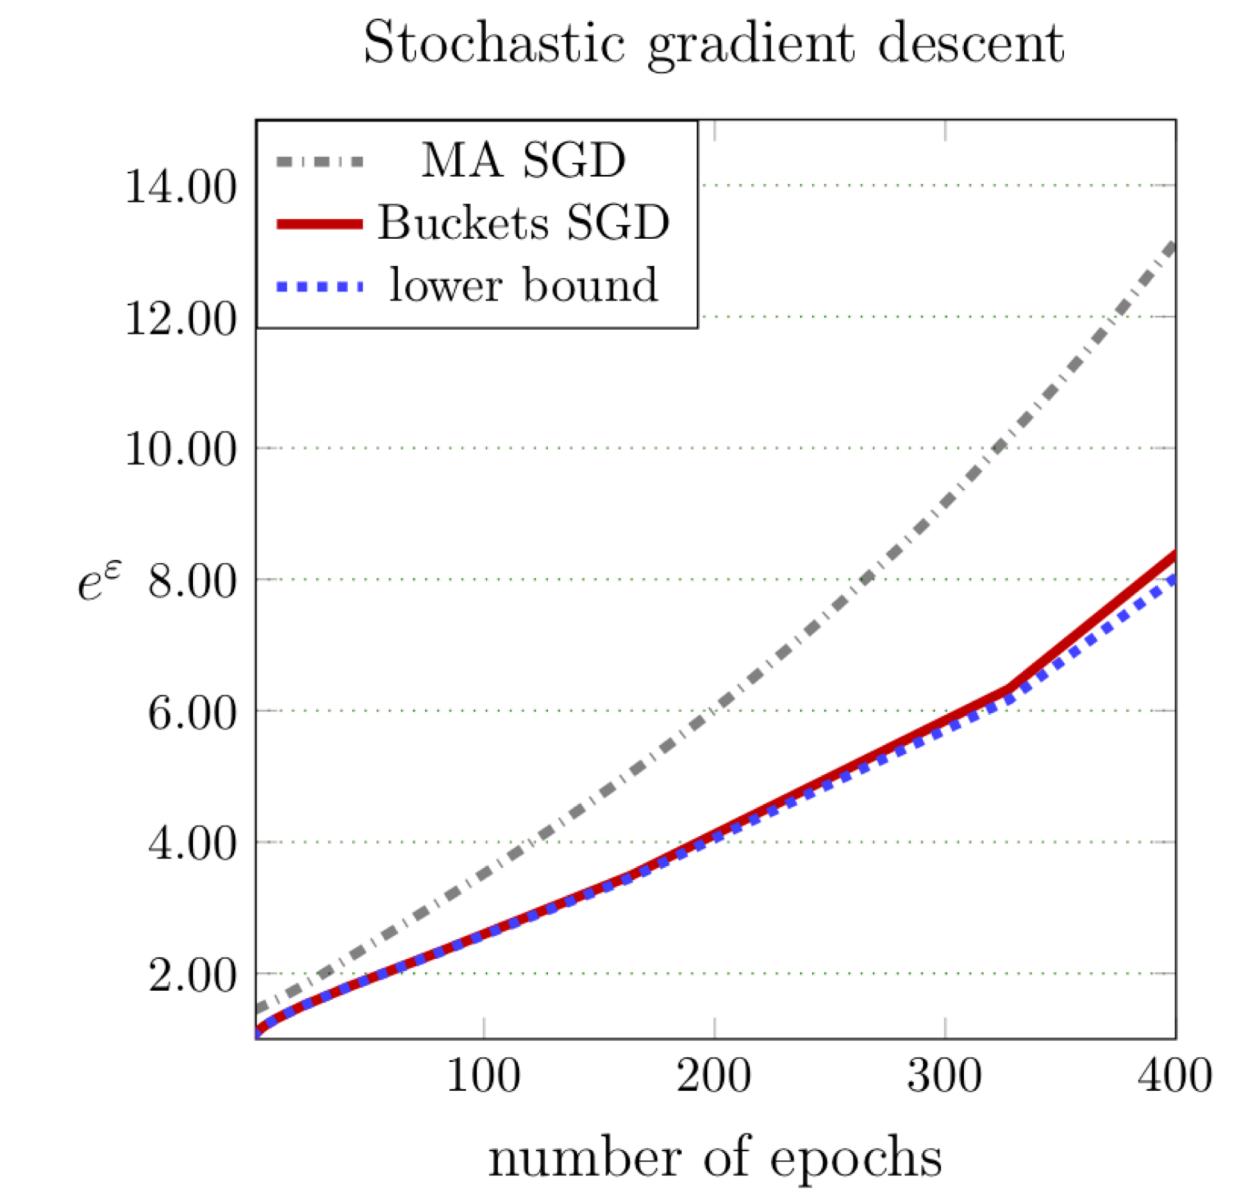
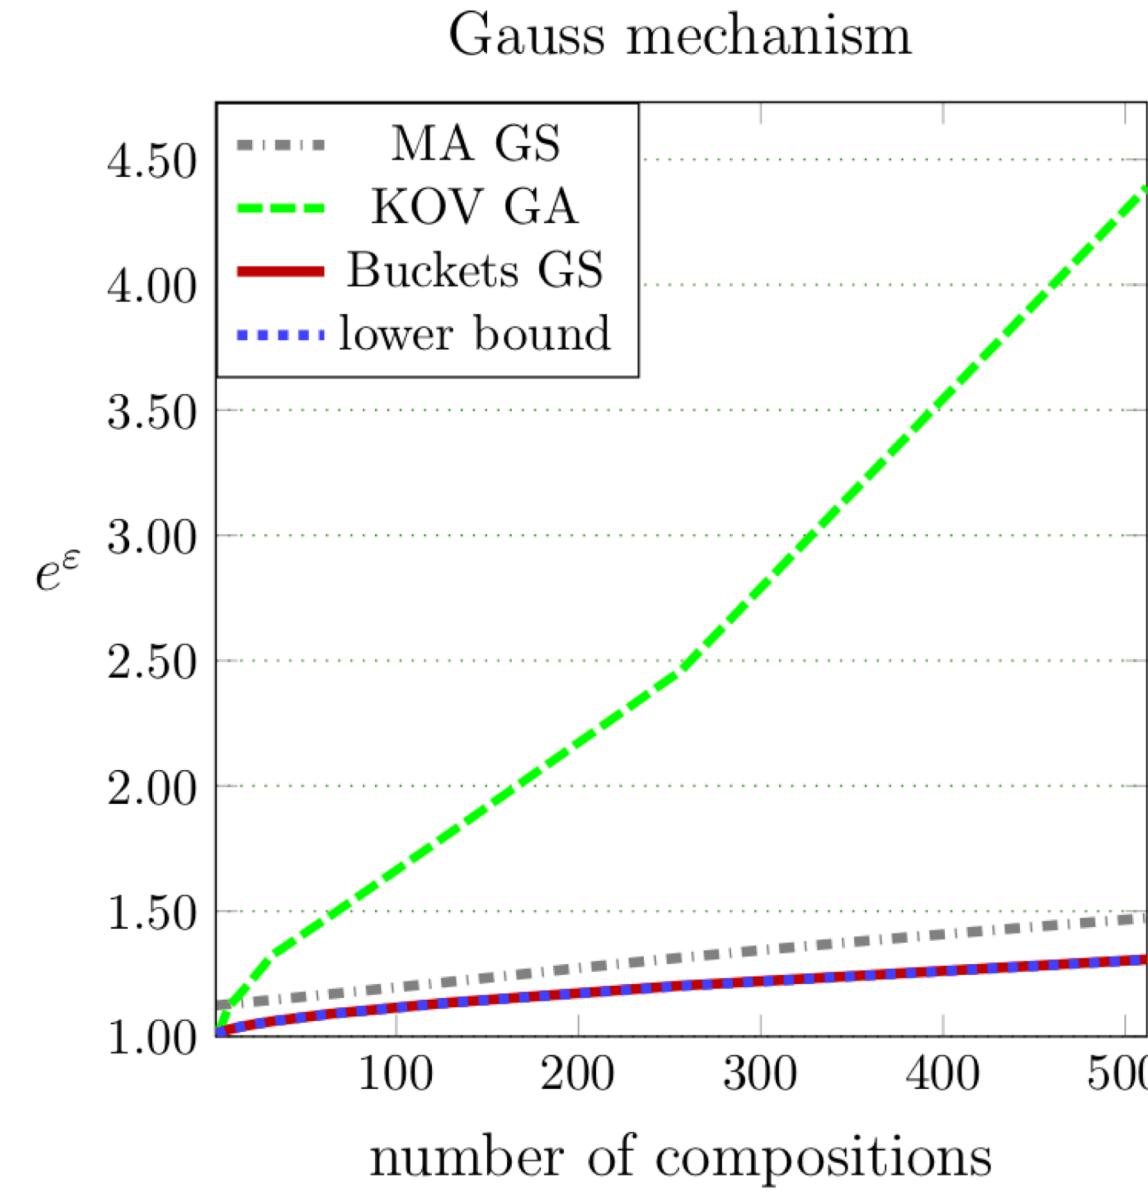
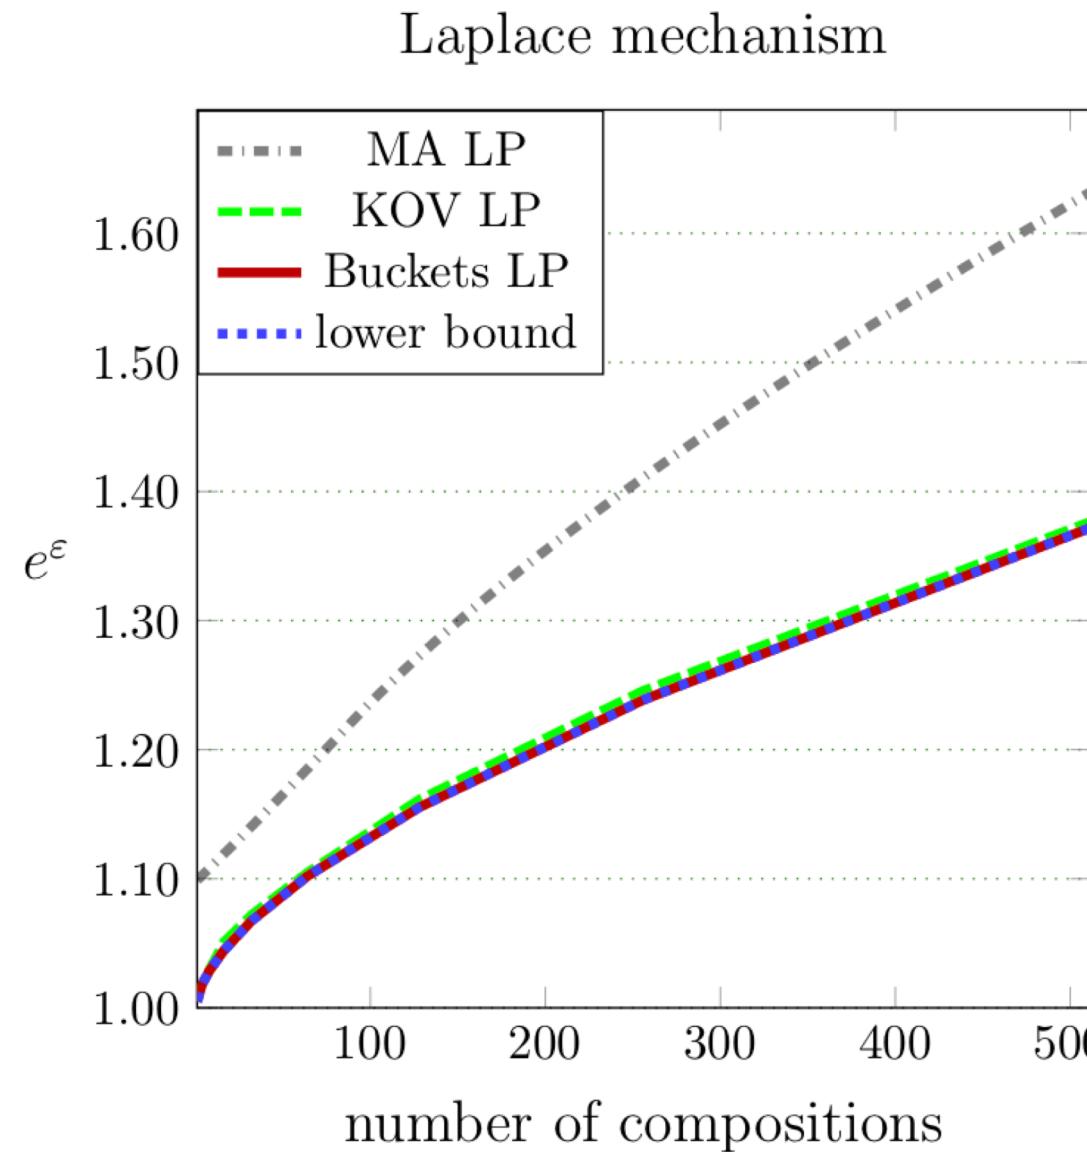
Comparison against the State-Of-The-Art

Laplace & Gauss after 512 observations, SGD after 60 000 observations
CDP bound not applicable to SGD (and CoverUp)



The rise of epsilon (over time)

We compute the minimal e^ε for which we can still achieve $\delta(\varepsilon) \leq 10^{-4}$ over the number of compositions.



Conclusion

Privacy Buckets: novel approximation of leakage

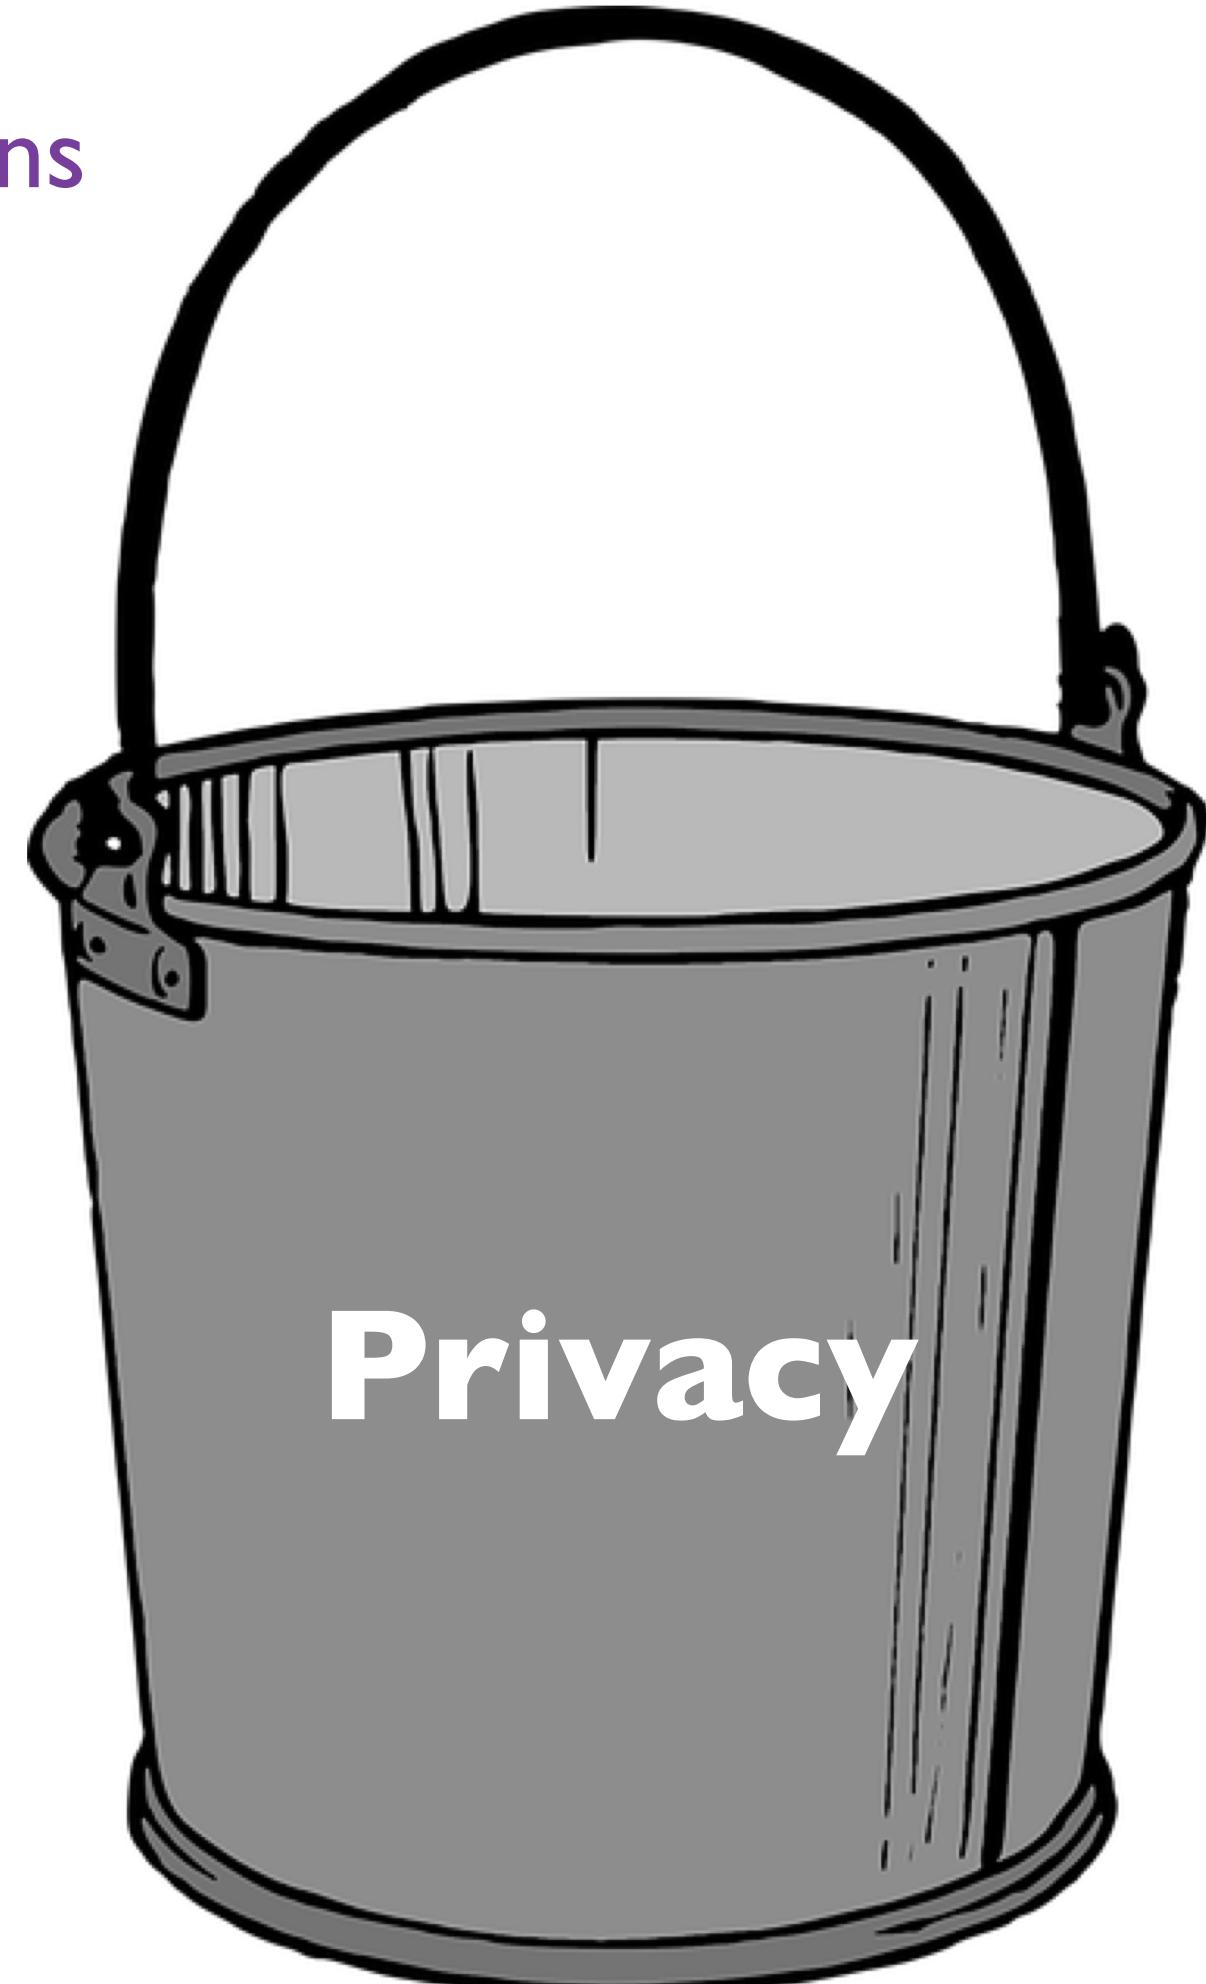
Computes **upper** and **lower** bounds for
Approximate Differential Privacy after **r** observations

Github repository:

<https://github.com/sommerda/privacybuckets>

To make our results more accessible:

- webpage to play around with (http for now)
<http://privacybuckets.space>



Thanks for your attention!
Questions?

privacybuckets.space